

Re-weighting each recording channel corrects most long-term decoder drift in people with implanted brain–computer interfaces

Md. Intiaj Alam Sajin^{1*}, Esm E Moula Chowdhury Abha¹ and Md Wahiduzzaman Suva¹

^{1*}Department of Computer Science, American International University-Bangladesh, Dhaka, 1229, Bangladesh.

*Corresponding author(s). E-mail(s): imtiajsajin@gmail.com;
Contributing authors: esmechowdhuryabha@gmail.com;
wshuvo360@gmail.com;

Abstract

Implanted brain–computer interfaces can restore communication and movement after paralysis, but their decoders lose accuracy as recordings change. Which changes matter, and how to correct them cheaply, remains unclear. Here, we show that in people, most of this loss can be undone by re-learning one weight per recording channel. Using one pipeline on public long-term recordings from three monkeys and 14 people, we found that decoder decay rates differed about a hundredfold between individuals, and weak regularization exaggerated decay. In all three people with decoding data, this re-weighting restored most of the accuracy lost over days to more than a year, for linear and recurrent decoders alike. In monkeys, the same correction recovered much less as gaps grew. Label-free alignment of neural activity added nothing beyond rescaling each channel. Brief electrode silences were mostly session-to-session fluctuation. These results point to fast, low-cost recalibration for clinical systems.

Keywords: brain–computer interface, intracortical recording, Utah array, neural drift, decoder recalibration, electrode stability

1 Introduction

2 Intracortical brain-computer interfaces (BCIs) record from small populations of neu-
3 rons and turn their activity into commands. They have allowed people with paralysis
4 to move robotic arms and computer cursors¹⁻⁴, to type by attempted handwriting⁵
5 and to speak through a decoded voice⁶⁻⁸. Recent work describes daily home use over
6 years⁹. For these devices to become everyday tools, their decoders must stay accurate
7 with little effort from the user.

8 The recorded signals are not stable. Within a day, signals shift as electrodes
9 move or neurons change their firing¹⁰. Across days, the activity seen by each elec-
10 trode changes¹¹⁻¹³. Over months and years, arrays gradually lose the ability to record
11 spikes¹⁴⁻¹⁷. Most systems therefore recalibrate their decoders often, which takes time
12 away from the user¹⁸. Many methods aim to reduce this cost. Decoders can be trained
13 on data from many days¹⁹ or updated during use^{18,20-22}. Neural activity can be
14 aligned to a reference day without new labels²³⁻²⁶, and monitoring scores have been
15 proposed to signal when recalibration is needed²⁷.

16 Choosing among these approaches requires knowing what actually changes in the
17 recordings. If the decoded signal moves to different channels, a decoder must be rebuilt.
18 If each channel keeps its tuning but changes in strength, a much smaller correction
19 may suffice. These possibilities have rarely been separated, and most comparisons used
20 one dataset, weeks of data and different baselines²⁸. On the hardware side, long-term
21 studies report falling array yield and impedance^{14,15,17,29}, but how single electrodes
22 come and go is less well described.

23 Public long-term datasets now allow these questions to be asked across individuals
24 with one method. They include three years of finger decoding in one monkey³⁰ and
25 reaching recordings from two monkeys in another laboratory^{24,31}. They also include
26 closed-loop cursor sessions spanning two to seven years in three BrainGate partici-
27 pants, array recordings from 20 arrays in 14 participants¹⁷, and closed-loop sessions
28 with fixed decoders²⁷.

29 Here, we analysed these datasets with a single pipeline. We trained decoders on
30 one day and tested them on later days after a ladder of increasingly powerful cor-
31 rections. Decay rates differed about a hundredfold between individuals, and weakly
32 regularized decoders exaggerated decay. In all three people with decoding data, re-
33 learning one weight per channel restored most of the lost accuracy at every gap tested,
34 for linear and recurrent decoders. In monkeys, this correction recovered much less at
35 long gaps, even when we decoded the same variable as in humans. Label-free align-
36 ment added nothing beyond rescaling each channel, and statistics of neural change
37 predicted decoder accuracy little better than elapsed time. At the electrode level,
38 brief silences were mostly session-to-session fluctuation, whereas silences lasting over
39 a month were real and often reversed. We release the code, a calibrated simulator and
40 all derived results.

Results

A common pipeline for six decoding datasets

We assembled four public datasets (Fig. 1a and Supplementary Table 1). Decoder analyses used six individuals. Monkey N was recorded over 1,242 days with Utah arrays in motor cortex during a two-finger task^{30,32}. Monkeys C and M were recorded over three and 1.5 years during reaching^{24,31}. BrainGate participants T6, T5 and T9 contributed 124, 92 and 117 closed-loop cursor sessions over 3.1, 7.2 and 2.0 years¹⁷. Electrode analyses used all 20 arrays from 14 BrainGate participants (2,289 sessions) and monkey N. Monitoring analyses used closed-loop sessions in which participants T11 and T5 controlled a cursor with a fixed decoder²⁷.

For each individual, we trained a linear decoder on one session and tested it on later sessions at gaps of 1 to 480 days. Before testing, every channel of the later session was z-scored with that session’s own mean and standard deviation. This step, which we call renormalization, needs no labels; the result is the renormalized fixed decoder. We then applied a ladder of corrections (Fig. 1b and Methods). Label-free rungs updated channel means or aligned neural activity to the training day. Supervised rungs used labelled trials from the later day. They either re-learned one weight per channel in front of the frozen decoder, or refitted the decoder while shrinking it toward the old one. A decoder trained on the later day served as the reference. We report retention, the ratio of a decoder’s accuracy (coefficient of determination, R^2) to that of the reference on the same test trials.

Decoders decay at rates that differ a hundredfold

Without any correction, decoders failed: reusing the training day’s normalization gave median R^2 below zero at most gaps (Supplementary Table 2). We therefore treat the renormalized fixed decoder as the baseline.

How much accuracy survives a single night differed between species (Fig. 2a,b). For true one-day gaps, retention was 0.74 in monkey N (95% confidence interval 0.62–0.83; $n = 21$ pairs), 0.70 in monkey C (0.58–0.77; $n = 31$) and 0.80 in monkey M (0.70–0.86; $n = 12$). In participant T9 it was 0.58 (0.33–0.76; $n = 19$). One-day pairs were rare in T6 ($n = 5$; retention 0.41, interval -2.45 – 0.99) and T5 ($n = 2$). At two days, retention was 0.79 in T6 ($n = 19$) and 0.59 in T5 ($n = 15$). Across the six individuals, retention after one to two days differed (Kruskal–Wallis test, $H = 11.4$, d.f. = 5, $P = 0.045$; one value per training session).

Longer-term decay differed far more. Retention fell to one half within 1.1 days in T9, 3.7 days in T5, 8.1 days in monkey N and 9.2 days in monkey C. It took 34 days in monkey M and 125 days in T6 (Fig. 2a,c). The reference decoder itself stayed equally accurate across gaps (Supplementary Table 2), so the decline reflects a growing mismatch between the decoder and the day, not poorer recordings.

Weak regularization exaggerates decay

Part of the measured decay came from the decoder rather than the recordings. With a ridge penalty of $\alpha = 10^4$ instead of the small values used in the original dataset




Fig. 1 Data and analysis design. **a**, Sessions used in this study, by individual and analysis, against years since each individual's first session. Each tick is one session. Blue, decoder drift (monkeys N, C and M; human participants T6, T5 and T9); green, electrode statistics (20 arrays in 14 BrainGate participants); orange, closed-loop monitoring (T11 and T5). Counts are per analysis; some human sessions enter more than one analysis. **b**, The correction ladder. A decoder trained on day i is tested on a later day j after increasingly powerful corrections: none; label-free corrections (rescaling each channel, aligning neural activity); and corrections that use labelled trials from day j (re-learning one weight per channel; refitting the decoder while shrinking it toward the old one). A decoder trained on day j is the reference.

papers³⁰, the time to half retention rose from 8.1 to 16 days in monkey N and from 125 to 229 days in T6 (Fig. 2c,d). It changed little in monkeys C (9.2 to 10 days) and M (34 to 37 days), whose original decoders were already more strongly regularized, and in T5 (3.7 to **??T5:thalfuned** days) and T9 (1.1 to **??T9:thalfuned** days). Per training session, absolute R^2 on later days rose by a median of 0.061 in monkey N (two-sided Wilcoxon signed-rank test, $P = 8.2 \times 10^{-15}$) and 0.019 in T6 ($P = 9.2 \times 10^{-22}$).

A simple rule captures this benefit using training-day data only: choose the largest penalty whose same-day accuracy stays within 1% of the best (Fig. 5a,b and Supplementary Table 6). Compared with $\alpha = 0.1$, it raised cross-day R^2 in every individual, with session-level median gains of 0.060 in monkey N (40 of 40 sessions; $P = 1.8 \times 10^{-12}$), 0.021 in T6 ($P = 1.8 \times 10^{-12}$), 0.033 in T5 ($P = 1.8 \times 10^{-12}$) and 0.045 in T9 ($P = 1.8 \times 10^{-12}$). Gains were small in monkeys C (0.017) and M (0.001). Same-day accuracy did not fall (median change 0.000 to +0.008). Ordinary same-day tuning chose similar penalties in most individuals (Methods), so the lesson is that small default penalties exaggerate drift, rather than that cross-validation fails.




Fig. 2 Decoders decay at rates that differ between individuals, and weak regularization exaggerates decay. **a**, Retention of the renormalized fixed decoder (decoding R^2 divided by that of a decoder trained on the test day) against days since training, with the ridge penalty used in the original dataset papers. Lines show medians and bands show 95% confidence intervals from a cluster bootstrap over training sessions. Circles and solid lines, monkeys; squares and dashed lines, humans. Blue, monkey N; orange, monkey C; green, monkey M; purple, T6; amber, T5; pink, T9. **b**, Retention after exactly one day (filled) and two days (open), with 95% confidence intervals; numbers give pairs. Markers are omitted when fewer than three pairs were available. **c**, Days until retention falls to one half with the default penalty (open) and with the tuned penalty (filled; logarithmic axis). **d**, As in **a**, for decoders trained with the tuned penalty ($\alpha = 10^4$). Pair counts are given in Supplementary Table 2.

Re-weighting each channel restores most accuracy in people

We next asked what kind of change the remaining loss reflects. Re-learning one weight per channel in front of the frozen decoder lets each channel's contribution grow, shrink or vanish, but keeps the decoder's pattern across time lags and outputs. In all three people, this correction restored most of the lost accuracy at every gap (Fig. 3a). It recovered 64–81% of the loss in T6, 72–85% in T5 and 76–84% in T9 (medians per gap; Fig. 3b). Retention after re-weighting stayed between 0.70 and 0.96 even 480 days after training. In the monkeys, the same correction recovered 14–34% of the loss in monkey N, 28–42% in C and 19–35% in M, and retention after re-weighting fell to 0.10–0.22 at 480 days.

The human result did not depend on the decoder. A recurrent network decoder (Methods), regularized to avoid memorizing the closed-loop sessions, was as accurate as the linear decoder in people (median R^2 0.24, 0.49 and 0.20 versus 0.25, 0.50 and 0.29). Re-weighting its input channels recovered 75–102%, 82–92% and 75–121% of its loss (Fig. 3b). In monkeys the network was more accurate than the linear decoder (median R^2 higher by 0.11, 0.33 and 0.42), but re-weighting again recovered less at long gaps (Supplementary Table 3).

Could the species difference reflect what was decoded? Human decoders predicted a two-dimensional direction, whereas monkey decoders predicted four kinematic signals. We therefore rebuilt the reaching monkeys' sessions to decode the direction of hand movement only, mirroring the human target (Methods). Re-weighting then recovered ??C:matchrec1% of the loss at one day and ??C:matchrec480% at 480 days in monkey C, and ??M:matchrec1% and ??M:matchrec480% in monkey M (Fig. 3c). The decoded variable therefore does not explain the difference.




Fig. 3 Re-weighting each channel restores most accuracy in people. **a**, Retention after re-learning one weight per channel in front of the frozen linear decoder, with up to 300 labelled trials from the test day. Lines, bands and colours as in Fig. 2a. **b**, Share of the loss recovered by re-weighting, (re-weighted R^2 minus renormalized R^2)/(reference R^2 minus renormalized R^2), for the linear decoder (filled) and a recurrent network decoder (open), median over all pairs with 95% cluster-bootstrap confidence intervals. **c**, Share of the loss recovered in the reaching monkeys when decoding hand kinematics (solid) or hand-movement direction only (dotted), the variable decoded in the human sessions, with the human median for comparison (grey). **d**, Retention against the number of labelled trials on the test day, pooled over the three human participants and all gaps, for re-weighting channels, refitting the decoder while shrinking it toward the old one, and training a fresh decoder. The dotted line marks the renormalized fixed decoder, which uses no labels.

How much labelled data does re-weighting need in people? Pooled over the three participants and all gaps, retention after re-weighting was `??hum:effL3n20` with 20 labelled trials, `??hum:effL3n50` with 50 and `??hum:effL3n300` with 300, against `??hum:effnone` without labels (Fig. 3d). A fresh decoder reached `??hum:effL6n20` with 20 trials and `??hum:effL6n50` with 50. A trial lasted about 1.6–3.9 s in these sessions, so 20 trials take under a minute in T5 and T6.

Rescaling each channel is the label-free step that matters

Several methods realign neural activity to a reference day without labels^{23–26}. Their benefits were usually measured against a fixed decoder or against one whose channel means alone were updated^{22,25}. We compared four label-free corrections with renormalization, which rescales each channel (Fig. 4a and Supplementary Table 4).

Rotating the dominant neural subspace onto the training day by orthogonal Procrustes analysis³³ did not help. Applied after updating channel means only, it lowered R^2 in monkey N by -0.063 (two-sided Wilcoxon signed-rank test over training sessions, $P = 3.4 \times 10^{-8}$). Applied after renormalization, it lowered R^2 by -0.019 ($P = 6.6 \times 10^{-9}$); in the humans the changes were `??T6:rotrendiff`, `??T5:rotrendiff` and `??T9:rotrendiff`. Renormalization itself improved on updating means only by 0.061 in monkey N ($P = 3.0 \times 10^{-11}$). Matching the full covariance of the channels (CORAL³⁴) changed R^2 by 0.011 in monkey N ($P = 1.9 \times 10^{-9}$), -0.000 in M, and -0.011 , -0.081 and `??T9:coraldiff` in the humans. A faithful implementation of the factor-analysis stabilizer of Degenhart and colleagues²³ improved on its own fixed latent decoder in none of the monkeys (change -0.008 in N, -0.090 in M) and did not consistently beat renormalization in the humans (changes -0.004 ,

0.125 and **??T9:stabrendiff**). That method was designed for closed-loop use with a latent decoder, and offline tests may understate it. In these data, however, the useful label-free step was to rescale each channel.

Neural change adds little beyond time since calibration

A monitor that warns when a decoder is failing could make recalibration less frequent. MINDFUL scores instability as the divergence between distributions of neural features, or of decoder outputs, on a test day and a reference day²⁷. Such scores correlate with decoder error, but both grow with time since calibration, so a correlation alone does not show that a score adds information.

Offline in monkey N (537 session pairs), elapsed time alone correlated with decoder accuracy (Spearman $|\rho| = 0.67$; Fig. 4b). Neural divergence correlated more weakly ($|\rho| = 0.34$), and its partial correlation fell to 0.16 after controlling for elapsed time. Output divergence behaved the same way ($|\rho| = 0.22$, partial 0.16). Adding both divergences to elapsed time slightly increased the cross-validated error of predicted R^2 (0.090 versus 0.085). A broader set of label-free statistics fared no better (Supplementary Note 1), although day-to-day performance beyond the time trend was reliable (split-half reliability 0.84).

We then re-analysed the public MINDFUL data, in which T11 and T5 controlled a cursor in closed loop with a fixed decoder (Fig. 4c and Supplementary Table 5). In T11 (17.1 days), neural divergence correlated with angular error ($r = 0.82$), consistent with the original report²⁷. Most of this association was shared with elapsed time (partial $r = 0.37$), and adding neural divergence did not improve leave-one-day-out prediction of error (17.5° versus 17.1°). Output divergence kept a partial correlation of 0.73, and adding it reduced the error to 11.9°. In T5, error was barely related to elapsed time over 32.9 days ($r = 0.17$), and both divergences tracked error (partial $r = 0.49$ for neural and 0.78 for output divergence). These comparisons involve few days, and output divergence and angular error are both computed from the cursor's motion, so part of their link is mechanical. Within these limits, elapsed time was as good a guide as neural statistics in most comparisons.

Shrinking toward the old decoder makes small recalibrations safe

When labelled data are scarce, refitting a decoder can make it worse. In monkey N (86 session pairs), we refitted the decoder by ridge regression shrunk toward the old decoder, penalizing changes in both the weights and the intercept^{20,21,35} (Fig. 5c and Supplementary Table 7). Choosing the shrinkage strength by cross-validation on a few trials was unreliable: with 10 trials, 10.5% of recalibrations ended worse than no recalibration (Fig. 5d). A fixed strong shrinkage, the value that worked best on pairs from other training sessions, lowered this to 1.2% (exact McNemar test on paired outcomes, $P = 0.008$; Holm-adjusted over five data budgets, $P = 0.039$). With 20 trials the rates were 8.1% and 0.0% ($P = 0.016$; adjusted $P = 0.062$). With 50 or more trials, and at gaps of 30 days or more, refitting from scratch did as well as the shrunk refit.




Fig. 4 Label-free corrections and monitoring add little to rescaling and elapsed time. **a**, Change in decoding R^2 relative to renormalization (rescaling each channel) for five label-free corrections, per individual, averaged over each training session’s pairs at gaps up to 120 days. Points show medians over training sessions and bars show 95% cluster-bootstrap confidence intervals. Colours and markers as in Fig. 2. **b**, Offline monitoring in monkey N (537 pairs): absolute Spearman correlation between decoder R^2 and the divergence of neural features, decoder outputs or both, before (light) and after (dark) controlling for elapsed days. The dashed line marks elapsed days alone. **c**, Closed-loop monitoring in the public MINDFUL data: leave-one-day-out mean absolute error in predicting angular error from elapsed days (grey), with neural divergence added (blue) or with output divergence added (orange).

Brief electrode silences are fluctuation, long ones are real

We next asked how single electrodes change over years. We used the BrainGate yield data, which report threshold-crossing rates and impedance for every electrode in every session¹⁷, and the matching data from monkey N. An electrode was active when its threshold-crossing rate exceeded 2 Hz, the definition used by BrainGate¹⁷. Array yield and impedance declined over years in both species (Fig. 6a,b). Impedance fell in 17 of 18 human arrays with measurements (median Spearman ρ with time -0.92) and in monkey N, from a median of 302 k Ω in the first year to 171 k Ω in the fourth ($\rho = -0.88$), replicating earlier reports^{14–17,29}.

Electrodes often fell silent and later recovered, but most such events were fluctuation rather than failure and repair (Fig. 6c and Supplementary Table 10). Of 730 initially active human electrodes that stayed below threshold for at least three consecutive sessions, 80% later recovered for at least three sessions. The same statistic computed after randomly shuffling each electrode’s session order, which destroys any temporal process, gave 79%. The match held when the threshold was fixed in microvolts for each electrode, so that changes in noise level could not move it (72% versus 72%), and when silence required rates below 0.5 Hz (71% versus 71%). Silences lasting at least 30 days were different. They were twice as common as in the shuffled data (168 versus 81), and 57% of them ended in recovery, more than in the shuffled data (41%; Fisher’s exact test, $P = 0.021$). Recovery here means the electrode again recorded spiking, not necessarily from the same neuron.

Spiking declined somewhat faster at the edge of the array than in its interior (Fig. 6d). Among initially active electrodes, the edge decline was steeper in 13 of 17 human arrays with at least 10 such electrodes (one-sided Wilcoxon signed-rank test with the array as the unit, $P = 0.036$), and decline correlated with distance from




Fig. 5 Simple recipes for robust decoders. **a**, Median decoding R^2 in monkey N against the ridge penalty α , on held-out trials of the training day (lightest) and on sessions 1, 7, 30, 120 and 480 days later (darker). The dotted line marks $\alpha = 0.1$, the value used in the original dataset paper. **b**, Gain in R^2 from the 1% rule (largest α whose same-day R^2 is within 1% of the best) over $\alpha = 0.1$, on the same day and at each gap, per individual; medians with 95% cluster-bootstrap confidence intervals. **c**, Retention in monkey N against the number of labelled trials on the test day, pooled over gaps, for refitting the decoder while shrinking it toward the old one, training a fresh decoder and re-weighting channels. **d**, Percentage of recalibrations (86 pairs) that ended more than 0.01 R^2 below the renormalized fixed decoder, with the shrinkage strength chosen by cross-validation on the available trials (grey) or fixed at the value that worked best on other training sessions (blue).

the array centre ($P = 0.013$). After adjusting for each electrode's initial firing rate, the edge effect weakened (12 of 17 arrays; $P = 0.087$). In monkey N, edge electrodes declined faster within both arrays for threshold crossings but not for spike-band power. These trends agree with models in which micromotion strains tissue most at the array edges³⁶, but they are modest and partly explained by initial activity.

A calibrated simulator predicts untested corrections

Methods for stable decoding are hard to compare without a common test bed²⁸. We built a feature-level simulator that takes a real session and returns it as it might appear days to years later (Methods). It combines mixing of signals between channels, gradual turnover of each channel's tuning and switching of electrodes between active and silent states. Switching rates came from monkey N's electrode statistics, and the mixing and turnover parameters were calibrated so that the simulated correction ladder matched monkey N's during its first 700 days (root-mean-square error 0.07 in retention; Fig. 7a).

We tested the simulator on the later days, which it had not seen, against two baselines (Fig. 7b). A simulator without drift was far worse (median absolute error 0.37–0.39). A stronger baseline simply reused the median retention measured in the first 700 days. For the corrections used in calibration, and for different amounts of labelled data, this empirical curve was as good as or better than the simulator (0.09 versus 0.12, and 0.07 versus 0.11). For corrections that the simulator was not fitted to, it was clearly better (0.06 versus 0.16). The simulator is therefore most useful for predicting how new corrections will behave. Real decoders decayed more slowly in the later years, and a calibration on that period gave weaker session-to-session mixing



Fig. 6 Brief electrode silences are fluctuation; long silences are real. **a**, Percentage of electrodes with spiking activity (threshold-crossing rate of at least 2 Hz) against years since implant, for 20 human arrays (purple) and the two arrays of monkey N (blue). **b**, Median impedance of each array relative to its first sessions (logarithmic axis). **c**, Percentage of silent episodes in initially active human electrodes that ended in recovery (at least three sessions above 4 Hz), observed (purple) and after shuffling each electrode's session order (grey), for four definitions of silence: at least three sessions below 2 Hz; the same with a fixed microvolt threshold per electrode; below 0.5 Hz; and lasting at least 30 days. Numbers give silent episodes. **d**, Difference in the decline of log threshold-crossing rate per year between edge and interior electrodes, per human array with at least 10 initially active electrodes, unadjusted (filled) and adjusted for initial firing rate (open). Negative values mean faster decline at the edge.

(0.77 versus 0.96) and slower turnover (time constant 527 versus 360 days), suggesting that drift slows as an implant ages in this monkey. Training decoders on simulated futures helped little (R^2 higher by 0.010) compared with stronger regularization (0.042) or training on several past sessions (0.056; Supplementary Table 12).

Discussion

We measured how decoders and electrodes change over years in three monkeys and 14 people with one pipeline. The central finding is that in people, most of the accuracy a decoder loses over days to more than a year is recovered by re-learning one weight per recording channel. This held in all three participants with decoding data, for linear and recurrent decoders, and with little labelled data. In monkeys, the same correction recovered much less at long gaps, even when we decoded the same variable.

This result has a simple interpretation. In these human recordings, the decoder's pattern across time lags and outputs remained valid for each channel, while the strength with which each channel carried the decoded signal changed. It agrees with the observation that tuning parameters on the same electrode tend to change together between days¹³. Our results do not reveal the cause. Changes in the number or proximity of neurons near each electrode would act this way^{37,38}, as would changes in noise. Why monkeys differ is also open. The human arrays were older at the start of recording, the human tasks were closed-loop, and the cortical areas and species differ. Each of these could change how tuning evolves.




Fig. 7 A calibrated simulator predicts untested corrections. **a**, Calibration on monkey N sessions before day 700. Points show real median retention and lines show the calibrated simulator, for the renormalized fixed decoder and after re-weighting channels. **b**, Validation on sessions after day 700: median absolute error between predicted and real retention for corrections used in calibration, corrections not used in calibration and amounts of labelled data not used in calibration. Dark blue, calibrated simulator; light blue, the median retention measured before day 700; grey, the simulator with drift switched off. **c**, Parameters calibrated separately on early (< 700 days) and late (≥ 700 days) sessions: session-to-session mixing, slow mixing, initial turnover ($\times 10$) and turnover time constant (years).

The practical implication is that recalibration in people can be small. Re-learning one weight per channel requires a short block of labelled trials, adds few parameters and keeps the decoder that the user has learned to control. It complements label-free methods, which in our data could not recover this kind of change beyond rescaling each channel. Three further recipes follow from our results. Daily per-channel renormalization should be the baseline against which new methods are compared. Ridge penalties should be tuned rather than left at small defaults, because weak regularization exaggerates apparent drift. And when recalibrating from a few trials, shrinking toward the previous decoder, including its intercept, avoids recalibrations that do more harm than good.

Decay rates varied about a hundredfold between individuals, from one day to several months to lose half of the accuracy. A fixed recalibration schedule will therefore be wrong for many users. The monitoring analyses suggest that statistics of neural change are a weak guide to when recalibration is needed, beyond elapsed time. Divergence of the decoder's output was more informative in closed loop, although part of that link is mechanical. Monitoring studies should report correlations after controlling for elapsed time and compare against a time-only baseline.

The electrode analyses refine common descriptions of array decline. Electrode activity near the yield threshold fluctuated strongly from session to session, and most short silences were no more than that. Yield counts from single sessions are therefore noisy, and decoders should not discard channels permanently after a few silent sessions. Silences lasting over a month were real and still reversed in more than half of cases. The edge effect we found is modest and partly explained by initial activity. Hahn and colleagues analysed yield, decoding signal quality and tuning stability in

278 the same human data¹⁷. We add electrode-level analyses of silence and recovery with
279 a shuffled null, and a test of edge-biased spiking loss with the array as the unit.

280 Our study has limitations. All analyses were offline. In closed loop, users adapt to
281 their decoder³⁹, which could slow or hide decay. Human decoding labels were inferred
282 from the direction from cursor to target, and the online decoder shaped both the
283 neural activity and the cursor path. Three participants and three monkeys cannot
284 establish population rates, and the datasets differ in task, signal type, array placement
285 and implant age. The simulator is calibrated to feature-level statistics in one monkey
286 and does not model biophysics. Finally, we tested label-free methods offline, where
287 they may perform differently than in closed loop.

288 Long-term public datasets made this study possible. We release our code, derived
289 results and the calibrated simulator so that others can test methods against realistic
290 drift. Our results suggest that durable clinical decoders may need less effort than
291 often assumed: a brief re-weighting of channels may keep a familiar decoder working
292 for months.

293 Methods

294 Datasets and ethics

295 All data were public and de-identified. The original studies obtained ethical approval
296 and informed consent, as described in the cited publications. We collected no new
297 data.

298 *Monkey N* (LINK; DANDI Archive 001201⁴⁰; ref.³⁰). One rhesus macaque had
299 Utah arrays in the hand area of primary motor cortex³². The release contains 312
300 sessions on 303 days over 1,242 days, in which the monkey moved two finger groups to
301 targets in centre-out or random-target blocks. We used the 96 released channels (64
302 on the medial array and 32 on the lateral array), with spike-band power and threshold
303 crossings at -4.5 times the root-mean-square noise in 20 ms bins⁴¹. Outputs were the
304 position and velocity of the two finger groups. Electrode impedance was stored on 228
305 days.

306 *Monkeys C and M* (DANDI Archive 000688⁴²; refs.^{24,31}). Both monkeys had Utah
307 arrays in motor and premotor cortex and performed centre-out and random-target
308 reaching. Monkey C had 68 sessions from October 2013 to October 2016, and monkey
309 M had 28 sessions from January 2014 to June 2015. We summed the spikes of all
310 sorted units on each electrode in 20 ms bins, giving 96 (C) or 192 (M) channels. Hand
311 position and velocity were interpolated to bin centres.

312 *BrainGate participants* (Dryad⁴³; ref.¹⁷). The yield data contain one recording
313 from each of 2,289 research sessions in 14 participants with 20 arrays (A1, S1–S3,
314 T1–T3 and T5–T11). For every electrode, they list threshold-crossing rates at several
315 thresholds (with the threshold in microvolts), impedance and grid position. Nineteen
316 arrays were in the hand area of motor cortex and one was in the middle frontal gyrus.
317 The decoding data contain closed-loop cursor sessions with spike-band power in 10 ms
318 bins, cursor and target positions and the periods of attempted movement. We used
319 all sessions of T6 (124 sessions over 3.1 years, one array, 96 channels), T5 (92 sessions
320 over 7.2 years, two arrays, 192 channels) and T9 (117 sessions over 2.0 years, two

arrays, 192 channels). Sessions contained a median of 367, 551 and 275 trials, with median movement periods of 1.6, 1.7 and 3.9 s.

MINDFUL (Dryad⁴⁴; ref.²⁷). Closed-loop cursor blocks in which T11 (15 days) and T5 (6 days) used a fixed decoder, with neural features in 20 ms bins (spike-band power for T11 and threshold crossings for T5), the decoder’s velocity output and the angular error of the cursor.

Features, decoders and accuracy

For monkey N and the human participants we used spike-band power (human 10 ms bins averaged in pairs). For monkeys C and M we used per-electrode spike counts. Each channel was z-scored with the mean and standard deviation of the decoding session’s training segment, computed without labels (renormalization). The linear decoder was ridge regression⁴⁵ on 8 causal lags (160 ms) of the z-scored features, with an unpenalized intercept. The default penalty was $\alpha = 0.1$ for monkey N and the humans, as in ref.³⁰, and $\alpha = 1$ for monkeys C and M, whose spike-count features are sparser. Outputs were finger position and velocity for monkey N, hand position and velocity for monkeys C and M, and for the humans the intended movement direction, defined as the unit vector from cursor to target during attempted movement. Accuracy was the coefficient of determination R^2 , averaged over outputs and computed on held-out trials of the test session. In monkey N the first 300 of 375 trials trained the decoder, as in ref.³⁰; in the other datasets the first 80% of trials did.

The recurrent decoder was a long short-term memory network⁴⁶ with 64 hidden units, input dropout⁴⁷ of 0.5 and 20-bin input windows (400 ms). It was trained with AdamW⁴⁸ (learning rate 10^{-3} , weight decay 0.01) for up to 4,000 iterations of 256 samples, with early stopping on the last 20% of training trials checked every 100 iterations. Without these measures, a 256-unit network memorized the human training sessions (training R^2 0.97, test R^2 -1.1 on one session). The network ladder used 25 training sessions per individual chosen at random with a fixed seed.

Session pairs

For each training session i and target gap $g \in \{1, 7, 30, 120, 480\}$ days, we selected the later session j of the same task type whose gap was closest to g , within 25% of g or one day, whichever was larger. The one-day set therefore contains one- and two-day gaps, which we report separately. In monkey N, 80 training sessions were drawn at random with a fixed seed; the full monkey N set (537 pairs) also includes further gaps and was used for monitoring. All human sessions were treated as one task type. Pair counts are listed in Supplementary Table 2.

Correction ladder

Let $\mathbf{x}_j(t) \in \mathbb{R}^C$ be the raw features of test day j , $\boldsymbol{\mu}_d$ and $\boldsymbol{\sigma}_d$ the channel means and standard deviations of day d , and f_i the decoder trained on day i .

- *No correction*: $f_i((\mathbf{x}_j - \boldsymbol{\mu}_i)/\boldsymbol{\sigma}_i)$.
- *Mean updating*: $f_i((\mathbf{x}_j - \boldsymbol{\mu}_j)/\boldsymbol{\sigma}_i)$.

- *Renormalization*: $f_i(\mathbf{z}_j)$ with $\mathbf{z}_j = (\mathbf{x}_j - \boldsymbol{\mu}_j)/\boldsymbol{\sigma}_j$.
- *Re-weighting*: f_i applied to $g_c z_{j,c}$ for each channel c , with one offset per output. The weights g_c were shrunk toward 1 and the offsets toward the old intercept; the penalty was chosen from $\{10^{-6}, \dots, 10^2\}$ by five-fold cross-validation over contiguous blocks of trials.
- *Ridge to prior*: a new decoder $(\mathbf{W}, \mathbf{b})$ fitted on the lagged features $\mathbf{H}$ of day j by minimizing

$$\|\mathbf{Y} - \mathbf{H}\mathbf{W} - \mathbf{1}\mathbf{b}^\top\|_F^2 + \lambda (\|\mathbf{W} - \mathbf{W}_i\|_F^2 + \|\mathbf{b} - \mathbf{b}_i\|_2^2), \quad (1)$$

where $(\mathbf{W}_i, \mathbf{b}_i)$ is the old decoder and λ was chosen from $\{0.1, 1, \dots, 10^6\}$ by cross-validation.

- *Reference*: a decoder trained on day j 's training trials.

Supervised rungs used the first n labelled trials of day j ($n = 300$, or all training trials if fewer, unless stated). Retention was the ratio of a rung's R^2 to the reference R^2 on the same test trials, and the share of the loss recovered by re-weighting was $(R_{\text{reweight}}^2 - R_{\text{renorm}}^2)/(R_{\text{ref}}^2 - R_{\text{renorm}}^2)$. The time to half retention was found by interpolating median retention linearly in log days. For the network decoder, re-weighting was a learned per-channel gain and offset in front of the frozen network, and the refit was fine-tuning with an L_2 penalty toward the old weights; penalties were chosen on the last 20% of labelled trials. We also computed a full input remap fitted by L-BFGS⁴⁹; because it can represent any linear decoder of the same form, we used it only as a calibration target for the simulator.

Matched-output control. For monkeys C and M, sessions were rebuilt with a two-dimensional unit-vector target, the direction of hand velocity, on bins where hand speed exceeded the session's 25th percentile. The ladder was run on the same pairs.

Data efficiency in humans. For 30 training sessions per participant (fixed seed), the ladder was run with $n \in \{10, 20, 50, 100, 300\}$ labelled trials.

Label-free corrections

Procrustes rotation. Let $\mathbf{V}_i, \mathbf{V}_j \in \mathbb{R}^{C \times k}$ hold the top $k = 16$ principal axes of the corrected training features on each day. The orthogonal matrix $\mathbf{R}$ minimizing $\|\mathbf{V}_j \mathbf{R} - \mathbf{V}_i\|_F$ was found by Procrustes analysis³³, and inputs were mapped as $(\mathbf{I} - \mathbf{V}_j \mathbf{V}_j^\top) \mathbf{z} + \mathbf{V}_i \mathbf{R}^\top \mathbf{V}_j^\top \mathbf{z}$. We applied it to renormalized features and, separately, to mean-updated features with $\mathbf{V}_j$ estimated from those features.

CORAL. Renormalized day- j features were re-coloured to match day i 's covariance, $\mathbf{z}' = \mathbf{C}_i^{1/2} \mathbf{C}_j^{-1/2} \mathbf{z}$ ³⁴, with covariances shrunk 5% toward a scaled identity.

Factor-analysis stabilizer. Following ref.²³, a factor-analysis model with 10 latent factors was fitted to each day's renormalized training features, and a ridge decoder on 8 lags of the day- i posterior latent means was trained. On day j , a model fitted without labels was aligned to day i by orthogonal Procrustes analysis of the loadings of stable channels, found by iteratively keeping the 60% of channels with the smallest loading residuals; latents were then inferred from stable channels only and passed to

the frozen latent decoder. The comparison baseline reused the day- i factor model on day j .

Monitoring analyses

In monkey N, for every pair we computed Gaussian Kullback–Leibler divergences between days i and j of the projections of renormalized features on day i ’s top ten principal axes (neural divergence) and of the decoder’s outputs (output divergence). We report raw Spearman correlations with R^2 , partial correlations given log elapsed days and time-blocked cross-validated linear prediction of R^2 . The additional statistics are described in Supplementary Note 1. Split-half reliability of the residual after the time trend used alternating pairs of trials and the Spearman–Brown formula. For the MINDFUL data we followed the original definitions²⁷: non-overlapping 60 s windows, divergences against the first day, and the median angular error over non-excluded trials (windows with fewer than 500 valid bins skipped). We report Pearson and partial correlations given elapsed days, and leave-one-day-out linear prediction of angular error; the reference day was excluded from evaluation.

Regularization and recalibration

Ridge decoders were trained with $\alpha \in \{0.1, 10, 10^3, 10^4, 3 \times 10^4, 10^5, 3 \times 10^5, 10^6\}$ on 40 training sessions per individual (22 for monkey M) and evaluated on held-out same-day trials and on later sessions at the five gaps. The 1% rule chose the largest α whose median same-day R^2 was at least 99% of the best median. We compared it with each training session’s own best same-day penalty: they did not differ in monkeys C and M or in T6, and the rule added 0.004 in monkey N ($P = 6.9 \times 10^{-5}$). The tuned-penalty ladder repeated the correction ladder on the same pairs with $\alpha = 10^4$ for both the day- i and the reference decoders. For recalibration with few trials we used 86 monkey N pairs and $n \in \{10, 20, 50, 100, 300\}$. The shrinkage λ in equation (1) was chosen per pair by cross-validation on the n trials, or as the value with the best median R^2 over pairs from other training sessions, which was 10^4 for every pair with up to 50 trials. A recalibration was harmful when its R^2 was more than 0.01 below that of the renormalized fixed decoder.

Electrode statistics

Statistics were computed per array. An electrode was active in a session when its threshold-crossing rate at -4.5 times the root-mean-square noise was at least 2 Hz. An initially active electrode (mean rate of at least 2 Hz over the first three sessions) was silent when it stayed below threshold for at least three consecutive observed sessions, and recovered when it later exceeded 4 Hz for at least three consecutive sessions. Missing sessions were skipped, not imputed. Control analyses (Supplementary Table 10) used thresholds of -3.5 and -5.5 times the noise; a fixed threshold per electrode in microvolts (its -4.5 value over the first three sessions, implemented by choosing in each session the noise multiple whose microvolt threshold was closest); removal of sessions with array-wide shifts (absolute median change in log rate above 1); a requirement for a mean waveform trough of at least 30 μ V; silence below 0.5 Hz;

441 and silences lasting at least 30 days. The null distribution randomly permuted each
 442 electrode’s session order. Switching rates were fitted by maximum likelihood to a two-
 443 state continuous-time Markov chain. For impedance, values of 4,000 k Ω or more were
 444 excluded and the trend was the Spearman correlation of the array median with day.
 445 For the edge effect, edge electrodes were those on the outer ring of the grid (10 $\times$ 10
 446 in humans and 8 $\times$ 8 in monkey N), each initially active electrode’s decline was the
 447 slope of $\log(\text{rate} + 0.1)$ against years, and arrays with at least 10 initially active elec-
 448 trodes were tested with the array as the unit. The adjusted analysis regressed slopes
 449 on edge membership and initial log rate within each array.

450 Simulator

451 The simulator transforms the z-scored features $\mathbf{Z} \in \mathbb{R}^{T \times C}$ of a real base session
 452 into those of a simulated session Δt days later (Supplementary Note 2). *Mixing*:
 453 $\mathbf{Z} \leftarrow \mathbf{Z}(\mathbf{I} + \mathbf{E})^\top$, with $\mathbf{E} = s(\Delta t)[(1 - \beta)\mathbf{G} + \beta\mathbf{L}]$, where $\mathbf{G}$ is a Gaussian matrix
 454 with zero diagonal, $\mathbf{L}$ a Gaussian matrix restricted to grid neighbours, $\beta = 0.5$ and
 455 $s(\Delta t)^2 = s_0^2 + [s_\infty(1 - e^{-\Delta t/\tau_s})]^2$. *Turnover*: each channel replaces a fraction ρ_c of its
 456 signal by a unit with rotated tuning in the session’s 20-dimensional latent space; ρ_c
 457 follows a beta distribution with mean $\rho_0 + (\rho_\infty - \rho_0)(1 - e^{-\Delta t/\tau_\rho})$ and concentration
 458 5. *Silencing*: each channel active in the base session becomes silent with probability
 459 $h_{\text{off}}/(h_{\text{off}} + h_{\text{on}})[1 - e^{-(h_{\text{off}} + h_{\text{on}})\Delta t}]$ and keeps 37% of its signal variance, the ratio
 460 of tuning strength of silent to active channels in monkey N. The switching rates
 461 ($h_{\text{off}} = 0.0038$ and $h_{\text{on}} = 0.0018$ per day) were measured from monkey N’s activity in
 462 the calibration period. The six mixing and turnover parameters were calibrated on 12
 463 base sessions before day 700 by minimizing the summed squared difference between
 464 simulated and real median retention for three rungs at five gaps, with the same correc-
 465 tion code and regularization as for real data, using a 24-point Sobol search⁵⁰ followed
 466 by Nelder–Mead⁵¹. Validation used 10 base sessions after day 700 and compared pre-
 467 dicted and real medians for all rungs and for $n \in \{20, 50, 100\}$. Augmentation added
 468 eight simulated futures per training session (Δt log-uniform between 1 and 120 days)
 469 and was compared with stronger regularization, random channel dropout with gain
 470 noise¹⁹ and training on the three previous sessions.

471 Statistics

472 Session pairs that share a training session are not independent. Confidence intervals
 473 for medians therefore came from a cluster bootstrap over training sessions⁵² (2,000
 474 resamples, percentile intervals), and paired tests used one value per training ses-
 475 sion (two-sided Wilcoxon signed-rank tests, over gaps up to 120 days for label-free
 476 corrections and over all gaps otherwise). Overnight retention was compared across
 477 individuals with a Kruskal–Wallis test. Harmful recalibrations were compared with
 478 exact McNemar tests on paired outcomes, Holm-adjusted over the five data budgets⁵³.
 479 Revival rates were compared with Fisher’s exact test. Electrode-edge tests used one
 480 value per array (one-sided sign and Wilcoxon signed-rank tests; the direction was pre-
 481 specified from earlier work³⁶). Exact P values are reported. Analyses used Python
 482 3.11 with NumPy, SciPy, pandas, scikit-learn and PyTorch.

Use of AI tools

An AI assistant (Claude, Anthropic) was used to write analysis code, run analyses and draft the text under the authors' direction. The authors reviewed the code, analyses and text and take full responsibility for the content.

Data availability

All data analysed in this study are public. LINK (monkey N) is available from the DANDI Archive at <https://doi.org/10.48324/dandi.001201/0.251023.2336>. The recordings of monkeys C and M are available at <https://doi.org/10.48324/dandi.000688/0.250122.1735>. The BrainGate yield and decoding data are available from Dryad at <https://doi.org/10.5061/dryad.x0k6djj1h>. The MINDFUL data are available from Dryad at <https://doi.org/10.5061/dryad.n2z34tn5s>. All derived results underlying the figures and tables are provided in the code repository.

Code availability

All code, including the simulator with its calibrated parameters, is available at <https://github.com/Imtiaj-Sajin/invasive-bci> under the MIT license. A versioned archive will be deposited on Zenodo upon publication.

Acknowledgements. We thank the participants of the BrainGate clinical trials and the teams who recorded and released the LINK, DANDI 000688, BrainGate and MINDFUL datasets. This work received no specific funding.

Author contributions. All authors contributed to the study design, analysis, interpretation and writing, and approved the final manuscript.

Competing interests. The authors declare no competing interests.

References

- [1] Hochberg, L. R. et al. Neuronal ensemble control of prosthetic devices by a human with tetraplegia. *Nature* **442**, 164–171 (2006).
- [2] Hochberg, L. R. et al. Reach and grasp by people with tetraplegia using a neurally controlled robotic arm. *Nature* **485**, 372–375 (2012).
- [3] Collinger, J. L. et al. High-performance neuroprosthetic control by an individual with tetraplegia. *Lancet* **381**, 557–564 (2013).
- [4] Pandarinath, C. et al. High performance communication by people with paralysis using an intracortical brain–computer interface. *eLife* **6**, e18554 (2017).
- [5] Willett, F. R., Avansino, D. T., Hochberg, L. R., Henderson, J. M. & Shenoy, K. V. High-performance brain-to-text communication via handwriting. *Nature* **593**, 249–254 (2021).

- 517 [6] Willett, F. R. et al. A high-performance speech neuroprosthesis. *Nature* **620**,
518 1031–1036 (2023).
- 519 [7] Card, N. S. et al. An accurate and rapidly calibrating speech neuroprosthesis. *N.*
520 *Engl. J. Med.* **391**, 609–618 (2024).
- 521 [8] Wairagkar, M. et al. An instantaneous voice-synthesis neuroprosthesis. *Nature*
522 **644**, 145–152 (2025).
- 523 [9] Card, N. S. et al. Long-term independent use of an intracortical brain–computer
524 interface for speech and cursor control. *Nat. Med.* **32**, 2504–2510 (2026).
- 525 [10] Perge, J. A. et al. Intra-day signal instabilities affect decoding performance in an
526 intracortical neural interface system. *J. Neural Eng.* **10**, 036004 (2013).
- 527 [11] Chestek, C. A. et al. Long-term stability of neural prosthetic control signals from
528 silicon cortical arrays in rhesus macaque motor cortex. *J. Neural Eng.* **8**, 045005
529 (2011).
- 530 [12] Downey, J. E., Schwed, N., Chase, S. M., Schwartz, A. B. & Collinger, J. L. Intra-
531 cortical recording stability in human brain–computer interface users. *J. Neural*
532 *Eng.* **15**, 046016 (2018).
- 533 [13] Bishop, W. E. et al. Self-recalibrating classifiers for intracortical brain–computer
534 interfaces. *J. Neural Eng.* **11**, 026001 (2014).
- 535 [14] Barrese, J. C. et al. Failure mode analysis of silicon-based intracortical micro-
536 electrode arrays in non-human primates. *J. Neural Eng.* **10**, 066014 (2013).
- 537 [15] Sponheim, C. et al. Longevity and reliability of chronic unit recordings using the
538 Utah, intracortical multi-electrode arrays. *J. Neural Eng.* **18**, 066044 (2021).
- 539 [16] Bjånes, D. A. et al. Quantifying physical degradation alongside recording
540 and stimulation performance of 980 intracortical microelectrodes chronically
541 implanted in three humans for 956–2130 days. *Acta Biomater.* **198**, 188–206
542 (2025).
- 543 [17] Hahn, N. V. et al. Performance of intracortical microelectrode arrays in people
544 with implanted brain–computer interfaces over a 20-year period. *Nat. Med.* (in
545 the press); preprint at <https://doi.org/10.1101/2025.07.02.25330310> (2025).
- 546 [18] Jarosiewicz, B. et al. Virtual typing by people with tetraplegia using a self-
547 calibrating intracortical brain–computer interface. *Sci. Transl. Med.* **7**, 313ra179
548 (2015).
- 549 [19] Sussillo, D., Stavisky, S. D., Kao, J. C., Ryu, S. I. & Shenoy, K. V. Making brain–
550 machine interfaces robust to future neural variability. *Nat. Commun.* **7**, 13749
551 (2016).

- 552 [20] Orsborn, A. L., Dangi, S., Moorman, H. G. & Carmena, J. M. Closed-loop
553 decoder adaptation on intermediate time-scales facilitates rapid BMI performance
554 improvements independent of decoder initialization conditions. *IEEE Trans.*
555 *Neural Syst. Rehabil. Eng.* **20**, 468–477 (2012).
- 556 [21] Dangi, S., Orsborn, A. L., Moorman, H. G. & Carmena, J. M. Design and analysis
557 of closed-loop decoder adaptation algorithms for brain–machine interfaces. *Neural*
558 *Comput.* **25**, 1693–1731 (2013).
- 559 [22] Wilson, G. H. et al. Long-term unsupervised recalibration of cursor-based intra-
560 cortical brain–computer interfaces using a hidden Markov model. *Nat. Biomed.*
561 *Eng.* **10**, 1466–1484 (2025).
- 562 [23] Degenhart, A. D. et al. Stabilization of a brain–computer interface via the align-
563 ment of low-dimensional spaces of neural activity. *Nat. Biomed. Eng.* **4**, 672–685
564 (2020).
- 565 [24] Gallego, J. Á., Perich, M. G., Chowdhury, R. H., Solla, S. A. & Miller, L. E. Long-
566 term stability of cortical population dynamics underlying consistent behavior.
567 *Nat. Neurosci.* **23**, 260–270 (2020).
- 568 [25] Ma, X. et al. Using adversarial networks to extend brain computer interface
569 decoding accuracy over time. *eLife* **12**, e84296 (2023).
- 570 [26] Karpowicz, B. M. et al. Stabilizing brain–computer interfaces through alignment
571 of latent dynamics. *Nat. Commun.* **16**, 4662 (2025).
- 572 [27] Pun, T. K. et al. Measuring instability in chronic human intracortical neural
573 recordings towards stable, long-term brain–computer interfaces. *Commun. Biol.*
574 **7**, 1363 (2024).
- 575 [28] Karpowicz, B. M. et al. Few-shot algorithms for consistent neural decoding
576 (FALCON) benchmark. *Adv. Neural Inf. Process. Syst.* **37** (2024).
- 577 [29] Chen, X. et al. Chronic stability of a neuroprosthesis comprising multiple adjacent
578 Utah arrays in monkeys. *J. Neural Eng.* **20**, 036039 (2023).
- 579 [30] Temmar, H. et al. Long-term Intracortical Neural activity and Kinematics
580 (LINK): an intracortical neural dataset for chronic brain–machine interfaces,
581 neuroscience, and machine learning. *Adv. Neural Inf. Process. Syst.* **38**, 154225–
582 154247 (2025).
- 583 [31] Perich, M. G., Gallego, J. Á. & Miller, L. E. A neural population mechanism for
584 rapid learning. *Neuron* **100**, 964–976.e7 (2018).
- 585 [32] Nason, S. R. et al. Real-time linear prediction of simultaneous and independent
586 movements of two finger groups using an intracortical brain–machine interface.

- 587 *Neuron* **109**, 3164–3177.e8 (2021).
- 588 [33] Schönemann, P. H. A generalized solution of the orthogonal Procrustes problem.
589 *Psychometrika* **31**, 1–10 (1966).
- 590 [34] Sun, B., Feng, J. & Saenko, K. Return of frustratingly easy domain adaptation. In
591 *Proc. 30th AAAI Conference on Artificial Intelligence* 2058–2065 (AAAI Press,
592 2016).
- 593 [35] Kuzborskij, I. & Orabona, F. Stability and hypothesis transfer learning. In *Proc.*
594 *30th International Conference on Machine Learning* 942–950 (PMLR, 2013).
- 595 [36] Forrest, A. M. et al. Finite element model predicts micromotion-induced strain
596 profiles that correlate with the functional performance of Utah arrays in humans
597 and non-human primates. *J. Neural Eng.* **22**, 066008 (2025).
- 598 [37] Polikov, V. S., Tresco, P. A. & Reichert, W. M. Response of brain tissue to
599 chronically implanted neural electrodes. *J. Neurosci. Methods* **148**, 1–18 (2005).
- 600 [38] Kozai, T. D. Y., Jaquins-Gerstl, A., Vazquez, A. L., Michael, A. C. & Cui,
601 X. T. Brain tissue responses to neural implants impact signal sensitivity and
602 intervention strategies. *ACS Chem. Neurosci.* **6**, 48–67 (2015).
- 603 [39] Gilja, V. et al. A high-performance neural prosthesis enabled by control algorithm
604 design. *Nat. Neurosci.* **15**, 1752–1757 (2012).
- 605 [40] Temmar, H. et al. LINK: Long-Term Intracortical Neural Activity and Kine-
606 matics. *DANDI Archive* <https://doi.org/10.48324/dandi.001201/0.251023.2336>
607 (2025).
- 608 [41] Nason, S. R. et al. A low-power band of neuronal spiking activity dominated
609 by local single units improves the performance of brain–machine interfaces. *Nat.*
610 *Biomed. Eng.* **4**, 973–983 (2020).
- 611 [42] Perich, M. G., Miller, L. E., Azabou, M. & Dyer, E. L. Long-term recordings of
612 motor and premotor cortical spiking activity during reaching in monkeys. *DANDI*
613 *Archive* <https://doi.org/10.48324/dandi.000688/0.250122.1735> (2025).
- 614 [43] Hahn, N. V. et al. Data from: Performance of intracortical microelectrode arrays
615 in people with implanted brain–computer interfaces over a 20-year period. *Dryad*
616 <https://doi.org/10.5061/dryad.x0k6djj1h> (2026).
- 617 [44] Pun, T. K. et al. Data from: Measuring instability in chronic human intracortical
618 neural recordings towards stable, long-term brain–computer interfaces. *Dryad*
619 <https://doi.org/10.5061/dryad.n2z34tn5s> (2024).
- 620 [45] Hoerl, A. E. & Kennard, R. W. Ridge regression: biased estimation for nonorthog-
621 onal problems. *Technometrics* **12**, 55–67 (1970).

- 622 [46] Hochreiter, S. & Schmidhuber, J. Long short-term memory. *Neural Comput.* **9**,
623 1735–1780 (1997).
- 624 [47] Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I. & Salakhutdinov, R.
625 Dropout: a simple way to prevent neural networks from overfitting. *J. Mach.*
626 *Learn. Res.* **15**, 1929–1958 (2014).
- 627 [48] Loshchilov, I. & Hutter, F. Decoupled weight decay regularization. In *Proc. 7th*
628 *International Conference on Learning Representations* (2019).
- 629 [49] Liu, D. C. & Nocedal, J. On the limited memory BFGS method for large scale
630 optimization. *Math. Program.* **45**, 503–528 (1989).
- 631 [50] Sobol', I. M. On the distribution of points in a cube and the approximate
632 evaluation of integrals. *USSR Comput. Math. Math. Phys.* **7**, 86–112 (1967).
- 633 [51] Nelder, J. A. & Mead, R. A simplex method for function minimization. *Comput.*
634 *J.* **7**, 308–313 (1965).
- 635 [52] Field, C. A. & Welsh, A. H. Bootstrapping clustered data. *J. R. Stat. Soc. B* **69**,
636 369–390 (2007).
- 637 [53] Holm, S. A simple sequentially rejective multiple test procedure. *Scand. J. Stat.*
638 **6**, 65–70 (1979).