

How intracortical brain–computer interfaces age and what keeps them working

Md. Intiaj Alam Sajin^{1*}, Esm E Moula Chowdhury Abha¹ and
Md Wahiduzzaman Suva¹

^{1*}Department of Computer Science, American International
University-Bangladesh, Dhaka, 1229, Bangladesh.

*Corresponding author(s). E-mail(s): imtiajsajin@gmail.com;
Contributing authors: esmechowdhuryabha@gmail.com;
wshuvo360@gmail.com;

Abstract

Implanted brain–computer interfaces can restore communication and movement after paralysis, but their decoders lose accuracy as recordings change over days and years. How fast this happens, and which corrections work, have not been measured consistently across individuals. Here, we analyzed public long-term recordings from three monkeys and 14 people with one pipeline. A fixed decoder lost half of its accuracy within one day to four months, depending on the individual, and weak regularization exaggerated this decay. In people, about 100 labeled trials restored most of the accuracy with any well-regularized refit, whereas label-free alignment added little beyond rescaling each channel. Corrections recovered far more when each channel kept its identity, and per-channel corrections tracked changes in each electrode’s firing. Brief electrode silences matched session-to-session fluctuation, but silences longer than a month were real and often reversed. These results provide evidence-based baselines for building durable interfaces.

Keywords: brain–computer interface, intracortical recording, Utah array, neural drift, decoder recalibration, electrode stability

¹ Introduction

² Intracortical brain–computer interfaces (BCIs) record from small populations of neu-
³ rons and turn their activity into commands. They have allowed people with paralysis

to move robotic arms and computer cursors¹⁻⁴, to type by attempted handwriting⁵ and to speak through a decoded voice⁶⁻⁸. Recent work describes daily home use over years⁹. For these devices to become everyday tools, their decoders must stay accurate with little effort from the user.

The recorded signals are not stable. Within a day, signals shift as electrodes move or neurons change their firing¹⁰. Across days, the activity seen by each electrode changes¹¹⁻¹³, and over months and years, arrays gradually lose the ability to record spikes¹⁴⁻¹⁷. Decoders therefore lose accuracy over time^{11,18-20}, and most systems recalibrate often, which takes time away from the user^{21,22}. Many methods aim to reduce this cost. Decoders can be trained on data from many days²³, updated during use^{20,21,24,25}, or given input layers learned for each day^{5,26}. Neural activity can be aligned to a reference day without new labels²⁷⁻³¹, and monitoring scores have been proposed to signal when recalibration is needed³².

Building on these approaches requires answers to three practical questions. How fast do decoders decay in different individuals, and how much of the measured decay comes from the decoder itself? Which corrections recover the loss, and with how much labeled data? And what happens at the electrodes over the years? Answers have come mostly from single datasets, spans of weeks and different baselines³³, which makes results hard to compare.

Public long-term datasets now allow these questions to be asked across individuals with one method. They include three years of finger decoding in one monkey¹⁹ and reaching recordings from two monkeys in another laboratory^{28,34}. They also include closed-loop cursor sessions spanning two to seven years in three BrainGate participants, recordings from 20 arrays in 14 participants¹⁷, and closed-loop sessions with fixed decoders³².

Here, we analyzed these datasets with a single pipeline. We trained decoders on one day, tested them on later days and applied a ladder of corrections. Decay rates differed about a hundredfold between individuals, and weak regularization exaggerated decay. In people, about 100 labeled trials restored most of the accuracy with any well-regularized refit, and label-free alignment added little beyond rescaling each channel. Corrections recovered far more when channels kept their identity, and the per-channel corrections tracked each electrode’s changing activity. At the electrode level, brief silences matched session-to-session fluctuation, whereas silences lasting over a month were real and often reversed. We release the code, a calibrated simulator and all derived results.

Results

A common pipeline for six decoding datasets

We assembled four public datasets (Fig. 1a and Supplementary Table 1). Decoder analyses used six individuals. Monkey N was recorded over 1,242 days with Utah arrays in motor cortex during a two-finger task^{19,35}. Monkeys C and M were recorded over three and 1.5 years during reaching^{28,34}. BrainGate participants T6, T5 and T9 contributed 124, 92 and 117 closed-loop cursor sessions over 3.1, 7.2 and 2.0 years¹⁷. Electrode analyses used all 20 arrays from 14 BrainGate participants (2,289 sessions)




Fig. 1 Data and analysis design. **a**, Sessions used in this study, by individual and analysis, against years since each individual's first session. Each tick is one session. Blue, decoder analyses (monkeys N, C and M; human participants T6, T5 and T9); green, electrode statistics (20 arrays in 14 BrainGate participants); orange, closed-loop monitoring (T11 and T5). Counts are per analysis; some human sessions enter more than one analysis. **b**, The correction ladder. A decoder trained on day i is tested on a later day j after increasingly powerful corrections: none; label-free corrections (renormalizing each channel; aligning neural activity); and corrections that use labeled trials from day j (re-weighting each channel; refitting the decoder while shrinking it toward the old one). A decoder trained on day j is the reference. Updating channel means only, a further label-free rung, is described in Methods.

and monkey N. Monitoring analyses used closed-loop sessions in which participants T11 and T5 controlled a cursor with a fixed decoder³².

For each individual, we trained a linear decoder on one session and tested it on later sessions at gaps of 1 to 480 days. Before testing, every channel of the later session was z-scored with that session's own mean and standard deviation. This step, renormalization, needs no labels and equalizes each channel's scale; the result is the renormalized fixed decoder. We then applied a ladder of corrections (Fig. 1b and Methods). Label-free corrections updated channel means or aligned neural activity to the training day. Corrections that use labeled trials from the later day either re-weighted each channel's contribution to the otherwise unchanged decoder, or refitted the decoder while shrinking it toward the old one. A decoder trained on the later day served as the reference. We report retention, the ratio of a decoder's accuracy (coefficient of determination, R^2) to that of the reference on the same test trials.




Fig. 2 Decoders decay at rates that differ between individuals, and weak regularization exaggerates decay. **a**, Retention of the renormalized fixed decoder (decoding R^2 divided by that of a reference decoder trained on the test day) against days since training, with the default ridge penalty. Lines show medians and bands show 95% confidence intervals from a cluster bootstrap over training sessions. Circles and solid lines, monkeys; squares and dashed lines, humans. Blue, monkey N; orange, monkey C; green, monkey M; purple, T6; amber, T5; pink, T9. **b**, Retention after exactly one day (filled) and two days (open), with 95% confidence intervals; numbers give pairs, and markers are omitted when fewer than three pairs were available. **c**, Days until median retention falls to one half with the default penalty (open) and with the tuned penalty (filled; logarithmic axis). **d**, As in **a**, for decoders trained with the tuned penalty ($\alpha = 10^4$). Pair counts are given in Supplementary Table 2.

Decoders decay at rates that differ a hundredfold

Without any correction, decoders failed: reusing the training day's normalization gave median R^2 below zero at most gaps (Supplementary Table 2). We therefore use the renormalized fixed decoder as the baseline.

How much accuracy survived a single night differed between individuals (Fig. 2a,b). For one-day gaps, retention was 0.74 in monkey N (95% confidence interval 0.62–0.83; $n = 21$ pairs), 0.70 in monkey C (0.58–0.77; $n = 31$), 0.80 in monkey M (0.70–0.86; $n = 12$) and 0.58 in participant T9 (0.33–0.76; $n = 19$). One-day pairs were rare in T6 ($n = 5$) and T5 ($n = 2$); at two days, retention was 0.79 in T6 ($n = 19$) and 0.59 in T5 ($n = 15$).

Longer-term decay differed far more. Retention fell to one half within 1.1 days in T9 (95% confidence interval 1.0–7.0), 3.7 days in T5 (1.0–9.5), 8.1 days in monkey N (3.7–10) and 9.2 days in monkey C (4.6–20). It took 34 days in monkey M (9.5–44) and 125 days in T6 (56–240; Fig. 2a,c). In monkey N and the three participants, the reference decoder was equally accurate at short and long gaps (Supplementary Table 2), so the decline reflects a growing mismatch between decoder and day rather than poorer recordings.

Weak regularization exaggerates decay

Part of the measured decay came from the decoder rather than the recordings. We had adopted the ridge penalty used in the original monkey N study¹⁹ ($\alpha = 0.1$; $\alpha = 1$ for the sparser spike counts of monkeys C and M). With a stronger penalty ($\alpha = 10^4$), the time to half retention roughly doubled in three individuals: monkey N (8.1 to 16 days), T6 (125 to 229 days) and T5 (3.7 to 7.9 days). It changed little in monkeys C

and M and in T9 (Fig. 2c,d). Per training session, absolute R^2 on later days rose by a median of 0.061 in monkey N (two-sided Wilcoxon signed-rank test, $P = 8.2 \times 10^{-15}$) and 0.019 in T6 ($P = 9.2 \times 10^{-22}$). Reports of decoder decay therefore depend on how strongly the decoder was regularized, and we use the tuned penalty for the remaining decoder analyses unless stated.

Rescaling each channel is the label-free step that matters

Several methods realign neural activity to a reference day without labels^{27–31}. Their benefits were usually measured against a fixed decoder or against one whose channel means alone were updated^{20,30}. We compared such corrections with renormalization (Fig. 3a and Supplementary Table 3).

Updating channel means alone was worse than renormalization in every individual; in monkey N the difference was 0.051 in R^2 (two-sided Wilcoxon signed-rank test over training sessions, $P = 1.7 \times 10^{-8}$). Rotating the dominant neural subspace onto the training day by orthogonal Procrustes analysis³⁶ did not close this gap. Applied after updating means, it lowered R^2 further in monkey N (by -0.068 ; $P = 1.3 \times 10^{-9}$). Applied after renormalization, it lowered R^2 slightly in every individual, for example by -0.015 in monkey N ($P = 4.4 \times 10^{-7}$). Matching the full channel covariance (CORAL³⁷) improved slightly on renormalization in monkey N (by 0.011; $P = 1.9 \times 10^{-9}$), did not differ in monkeys C and M, and was worse in the three participants. A faithful implementation of the factor-analysis stabilizer of Degenhart and colleagues²⁷ was better than renormalization in T5 (by 0.070; $P = 4.8 \times 10^{-9}$), did not differ in monkey N and was worse in the other four individuals. That method was designed for closed-loop use with a latent decoder, and offline tests may understate it. Methods that align latent dynamics with labeled data²⁸ or adversarial networks^{29–31} were not tested. Among the corrections we tested, however, the label-free step that mattered was rescaling each channel, and gains reported for alignment over mean-only baselines²⁰ may largely reflect this step.

Neural change adds little beyond time since calibration

A monitor that warns when a decoder is failing could make recalibration less frequent. MINDFUL scores instability as the divergence between distributions of neural features, or of decoder outputs, on a test day and a reference day³². Such scores correlate with decoder error, but both grow with time since calibration.

Offline in monkey N (537 session pairs), elapsed time alone correlated with decoder accuracy (Spearman $|\rho| = 0.67$; Fig. 3b). Neural divergence correlated more weakly ($|\rho| = 0.34$), and its partial correlation fell to 0.16 after controlling for elapsed time; output divergence behaved the same way. Adding both divergences to elapsed time slightly increased the cross-validated error of predicted R^2 (0.090 versus 0.085). In a pilot on the first year of monkey N, a broader set of label-free statistics also failed to improve on elapsed time, although day-to-day performance beyond the time trend was reliable (split-half reliability 0.84; Supplementary Note 1).

We then re-analyzed the public MINDFUL data, in which T11 and T5 controlled a cursor in closed loop with a fixed decoder (Fig. 3c and Supplementary Table 4). In




Fig. 3 Label-free corrections and monitoring add little to renormalization and elapsed time. **a**, Change in decoding R^2 relative to renormalization for five label-free corrections, per individual, over each training session’s pairs at all gaps. Points show medians over training sessions and bars show 95% cluster-bootstrap confidence intervals; values below the axis are printed. Colours and markers as in Fig. 2. **b**, Offline monitoring in monkey N (537 pairs): absolute Spearman correlation between decoder R^2 and the divergence of neural features, decoder outputs or both, before (light) and after (dark) controlling for elapsed days. The dashed line marks elapsed days alone. **c**, Closed-loop monitoring in the public MINDFUL data: leave-one-day-out mean absolute error in predicting angular error from elapsed days (grey), with neural divergence added (blue) or with output divergence added (orange).

T11 (15 days), neural divergence correlated with angular error across 60-s windows ($r = 0.82$), consistent with the original report³². Most of this association was shared with elapsed time (partial $r = 0.37$), and adding neural divergence did not improve leave-one-day-out prediction of error (17.5° versus 17.1°). Output divergence kept a partial correlation of 0.73, and adding it reduced the error to 11.9° . These window-level statistics are nested within few days, and output divergence and angular error are both computed from the cursor’s motion, so part of their link is mechanical. The six days available for T5 allow only a descriptive comparison. Within these limits, elapsed time was as good a guide as neural statistics.

Corrections recover more when channels keep their identity

Label-free corrections left much of the loss in place. We next asked what kind of change the remaining loss reflects, using re-weighting: one weight per channel in front of the otherwise unchanged decoder, fitted with labeled trials from the later day. Re-weighting lets each channel’s contribution grow, shrink or vanish, but keeps the decoder’s pattern across time lags and outputs. In the three participants it restored retention to 0.70–0.93 at every gap up to 480 days, recovering 65–74%, 69–83% and 73–80% of the loss in T6, T5 and T9 (Fig. 4a). In the monkeys it recovered –5–21%, 24–39% and 19–35%.

This difference between species mostly reflects what was decoded, not how the recordings changed. A two-dimensional target, such as the cursor-to-target direction decoded for the participants, can be fitted from many channel weights even when the decoder’s per-channel patterns are wrong. To test this, we re-weighted decoders whose channels had been randomly scrambled (Fig. 4b and Supplementary Table 5). In T5,

the scrambled decoder reached a retention of 0.77, close to the 0.83 of the correctly assigned decoder; in T9 the values were 0.55 and 0.78. In the monkeys, whose decoders predicted four kinematic signals, scrambled decoders could not be rescued (retention -0.19 , 0.02 and 0.02). Decoding hand-movement direction in monkeys C and M raised the share of the loss recovered to 36–60% and 30–48% (Supplementary Table 7), but did not reproduce the human setting.

Channel identity nevertheless mattered in every individual. Correctly assigned decoders always beat scrambled ones, and the difference was largest when weights could not change sign, so that scrambled channels could not be reversed. With non-negative weights, retention was 0.87 versus 0.36 in T6, 0.72 versus 0.57 in T5 and 0.72 versus 0.29 in T9, and 0.49 versus -0.17 , 0.57 versus 0.00 and 0.67 versus -0.00 in the monkeys. Constraining the weights to be non-negative cost little with correct channels. Drift therefore preserves much of each channel’s tuning pattern.

The per-channel weights also related to measured changes at each electrode (Fig. 4c and Supplementary Table 5). For each pair, we compared the fitted weights with the change in each electrode’s threshold-crossing rate between the two days, recorded independently in the BrainGate yield data. Channels whose spiking rose received larger weights in all three participants. The median Spearman correlation per training session was 0.04 in T6 (positive in 69% of sessions; two-sided Wilcoxon signed-rank test, $P = 0.017$), 0.03 in T5 (77%; $P = 8.9 \times 10^{-4}$) and 0.11 in T9 (92%; $P = 3.2 \times 10^{-10}$), against near zero after shuffling channels. In T9, channels whose impedance rose received smaller weights (median correlation -0.08 ; $P = 1.3 \times 10^{-7}$). These relations are weak for single channels but consistent, linking the decoder’s corrections to physical changes at the array.

The results held for nonlinear decoders. A recurrent network decoder (Methods), regularized to avoid memorizing the closed-loop sessions, was slightly less accurate than the linear decoder in people (median R^2 0.24, 0.49 and 0.20 versus 0.25, 0.50 and 0.31) and more accurate in monkeys. Re-weighting its input channels recovered 75–102%, 82–92% and 83–105% of its loss in the participants and 36–63%, 46–68% and 48–69% in the monkeys (Fig. 4d and Supplementary Table 6). Finally, allowing one scale per channel reduced the angle between decoders fitted on different days by 9–34° in every individual (Supplementary Table 8), so angles that allow only a global scale overstate how much decoders rotate across days²⁰.

A few minutes of labeled data restore decoders in people

How much labeled data does recalibration need? In the three participants, retention pooled over all gaps rose from 0.26 without labels to 0.73 with 50 labeled trials and 0.85 with 100 for a freshly trained, strongly regularized decoder (Fig. 5c and Supplementary Table 10). Refitting while shrinking toward the old decoder reached 0.67 and 0.77, and re-weighting reached 0.58 and 0.68. Movement periods lasted 1.6–3.9 s in these sessions, so 100 trials take a few minutes. Recovery was slower in T9 than in T5 and T6 (Supplementary Table 10). Re-weighting therefore indicates what changes, but it is not the most data-efficient recalibration in these sessions.

Two further choices made recalibration more reliable. First, the ridge penalty can be tuned on training-day data alone. We chose the largest penalty whose same-day




Fig. 4 Corrections recover more when channels keep their identity. **a**, Retention after re-weighting each channel in front of the frozen linear decoder (tuned penalty), with up to 300 labeled trials from the test day. Lines, bands, colours and markers as in Fig. 2a. **b**, Median retention after re-weighting the correctly assigned decoder (circles or squares) or a decoder with scrambled channels (crosses), with signed weights (upper row of each pair, filled) or non-negative weights (lower row, open); the vertical tick marks the renormalized fixed decoder. **c**, Spearman correlation between fitted channel weights and the change in each electrode's threshold-crossing rate between the two days, one point per training session (coloured), and the same after shuffling channels (grey); bars show medians. **d**, Share of the loss recovered by re-weighting, $(\text{re-weighted } R^2 \text{ minus renormalized } R^2) / (\text{reference } R^2 \text{ minus renormalized } R^2)$, for the linear decoder (filled) and a recurrent network decoder (open), with 95% cluster-bootstrap confidence intervals.

accuracy stayed within 1% of the best. This rule raised cross-day R^2 in every individual without lowering same-day accuracy (Fig. 5a,b and Supplementary Table 9). Session-level median gains were 0.060 in monkey N ($P = 1.8 \times 10^{-12}$) and 0.021, 0.033 and 0.045 in the participants. Ordinary same-day tuning usually chose a similar penalty (Methods). Second, with very few trials, refitting can make a decoder worse. In monkey N (86 session pairs), choosing the shrinkage strength by cross-validation on 10 trials made 10.5% of recalibrations worse than no recalibration (Fig. 5d). A fixed strong shrinkage, the value that worked best on other training sessions, lowered this to 1.2% (exact McNemar test, $P = 0.008$; Holm-adjusted over five data budgets, $P = 0.039$); with 20 trials the rates were 8.1% and 0.0% ($P = 0.016$; adjusted $P = 0.062$).

Brief electrode silences are fluctuation, long ones are real

We next asked how single electrodes change over years, using the BrainGate yield data, which report threshold-crossing rates and impedance for every electrode in every session¹⁷, and the matching data from monkey N. An electrode was active when its threshold-crossing rate exceeded 2 Hz, the definition used by BrainGate¹⁷. Array yield and impedance declined over years in both species (Fig. 6a,b). Impedance fell in 17 of 18 human arrays with measurements (median Spearman ρ with time -0.92) and in monkey N, from a median of 302 k Ω in the first year to 171 k Ω in the fourth ($\rho = -0.88$), replicating earlier reports^{14–17,38}.

Electrodes often fell silent and later recovered^{12,14}, but most such events were fluctuation (Fig. 6c and Supplementary Table 11). Of 730 initially active human electrodes that stayed below threshold for at least three consecutive sessions, 80% later




Fig. 5 Simple recipes for recalibration. **a**, Median decoding R^2 in monkey N against the ridge penalty α , on held-out trials of the training day (lightest) and on sessions 1, 7, 30, 120 and 480 days later (darker); the dotted line marks $\alpha = 0.1$. **b**, Gain in R^2 from the 1% rule (largest α whose same-day R^2 is within 1% of the best) over $\alpha = 0.1$, on the same day and at each gap, per individual; medians with 95% cluster-bootstrap confidence intervals. **c**, Retention in the three participants against the number of labeled trials on the test day, pooled over all gaps, for a decoder retrained on those trials, refitting while shrinking toward the old decoder, and re-weighting channels; the dotted line marks the renormalized fixed decoder. **d**, Percentage of recalibrations in monkey N (86 pairs) that ended more than 0.01 R^2 below the renormalized fixed decoder, with the shrinkage strength chosen by cross-validation on the available trials (grey) or fixed at the value that worked best on other training sessions (blue).

recovered for at least three sessions. When each electrode's session order was shuffled 200 times, which destroys any temporal process, the recovery rate was 78% (95% range 75–81%; permutation $P = 0.075$). The match held with a fixed threshold in microvolts for each electrode, so that changes in noise level could not move it, and when silence required rates below 0.5 Hz. Silences lasting at least 30 days were different. They were twice as common as in the shuffled data (168 versus 79 on average), and 57% of them ended in recovery, against 44% in the shuffled data (95% range 33–54%; $P = 0.010$). Recovery here means that the electrode again recorded spiking, not necessarily from the same neuron.

Spiking declined somewhat faster at the edge of the array than in its interior (Fig. 6d). Among initially active electrodes, the edge decline was steeper in 13 of 17 human arrays with at least 10 such electrodes (Wilcoxon signed-rank test with the array as the unit, one-sided $P = 0.036$). After adjusting for each electrode's initial firing rate, the effect weakened (12 of 17 arrays; one-sided $P = 0.087$). In monkey N, edge electrodes declined faster within both arrays for threshold crossings but not for spike-band power. These trends agree with models in which micromotion strains tissue most at the array edges³⁹, but they are modest and partly explained by initial activity.

A simulator predicts corrections it was not fitted to

Methods for stable decoding are hard to compare without a common test bed³³. We built a feature-level simulator that takes a real session and returns it as it might appear days to years later (Methods). It combines mixing of signals between channels,




Fig. 6 Brief electrode silences are fluctuation; long silences exceed a shuffled null. **a**, Percentage of electrodes with spiking activity (threshold-crossing rate of at least 2 Hz) against years since implant, for 20 human arrays (purple) and the two arrays of monkey N (blue). **b**, Median impedance of each array relative to its first sessions (logarithmic axis). **c**, Percentage of silenced initially active human electrodes that recovered (at least three sessions above 4 Hz), observed (purple) and after shuffling each electrode's session order 200 times (grey; bars show the 95% range), for four definitions of silence: at least three sessions below 2 Hz; the same with a fixed microvolt threshold per electrode; below 0.5 Hz; and lasting at least 30 days. Numbers give silenced electrodes. **d**, Difference in the decline of log threshold-crossing rate per year between edge and interior electrodes, per human array with at least 10 initially active electrodes, unadjusted (filled) and adjusted for initial firing rate (open). Negative values mean faster decline at the edge.

gradual turnover of each channel's tuning and switching of electrodes between active and silent states. Switching rates came from monkey N's electrode statistics, and the mixing and turnover parameters were calibrated so that the simulated correction ladder matched monkey N's during its first 700 days (root-mean-square error 0.07 in retention over the three fitted corrections and five gaps; Fig. 7a).

We tested the simulator on the later days, which it had not seen, against two baselines (Fig. 7b). A simulator without drift was far worse (median absolute error 0.27–0.39). A stronger baseline simply reused the median retention measured in the first 700 days. For the corrections used in calibration, and for different amounts of labeled data, this empirical curve was as good as or better than the simulator (0.09 versus 0.12, and 0.07 versus 0.11). For corrections that the simulator was not fitted to, it was clearly better (0.06 versus 0.16). The simulator is therefore most useful for predicting how new corrections will behave. Real decoders decayed more slowly in the later years. A calibration on that period gave weaker session-to-session mixing (0.77 versus 0.96) and slower turnover (time constant 527 versus 360 days; Fig. 7c and Supplementary Table 12). Drift may therefore slow as an implant ages, at least in this monkey. Training decoders on simulated futures helped little (R^2 higher by 0.010) compared with stronger regularization (0.042) or training on several past sessions (0.056; Supplementary Table 13).




Fig. 7 A calibrated simulator predicts corrections it was not fitted to. **a**, Calibration on monkey N sessions before day 700. Points show real median retention and lines show the calibrated simulator, for the renormalized fixed decoder and after re-weighting channels. **b**, Validation on sessions after day 700: median absolute error between predicted and real retention for corrections used in calibration, corrections not used in calibration and amounts of labeled data not used in calibration. Dark blue, calibrated simulator; light blue, the median retention measured before day 700; grey, the simulator with drift switched off. **c**, Parameters calibrated separately on early (< 700 days) and late (≥ 700 days) sessions: session-to-session mixing, slow mixing, initial turnover ($\times 10$) and turnover time constant (years).

Discussion

We measured how intracortical BCIs age, at the level of decoders and of electrodes, in three monkeys and 14 people with one pipeline. Five findings stand out. Decay rates differ about a hundredfold between individuals. Measured decay depends on how strongly the decoder is regularized. In people, a few minutes of labeled data restore most of the accuracy, whereas the label-free corrections we tested add little beyond rescaling each channel. Corrections work far better when each channel keeps its identity, and they track changes at each electrode. And most brief electrode silences are noise, while long ones are real and often reverse.

These findings have direct implications for building durable interfaces. Because decay rates vary so widely, a fixed recalibration schedule is unlikely to suit all users, and schedules may need to be set per person. Because weak regularization exaggerates decay, comparisons of decoder longevity are meaningful only between well-regularized decoders. Per-channel renormalization is the appropriate baseline for evaluating label-free methods; earlier comparisons against fixed or mean-only decoders^{20,30} probably overstated their benefit. For recalibration in people, any well-regularized refit with about 100 labeled trials recovered most of the accuracy in our data. When only a few trials are available, shrinking toward the previous decoder, including its intercept, prevents harmful updates. These are simple, inexpensive steps, and they set the bar that more sophisticated methods should clear.

The channel analyses clarify what changes. Earlier work modeled per-electrode drift as shifts in baseline over weeks^{13,21}, and human decoders absorb day-to-day changes with input layers learned for each day^{5,26}. Other studies describe changes in the preferred directions of individual features³² and rotation of decoder weights²⁰. In our data, much of each channel's tuning pattern survived for months:

decoders with correctly assigned channels consistently beat scrambled ones, and the per-channel corrections tracked each electrode’s firing. At the same time, our scrambled-channel control shows that high recovery from a per-channel correction is not, by itself, evidence of per-channel drift when the decoded variable is low-dimensional. Correction-based analyses of drift, including ours, need such controls. Changes in the number or proximity of neurons near each electrode would produce per-channel changes^{40,41}, but our data cannot identify the cause.

The electrode analyses refine common descriptions of array decline. Activity near the yield threshold fluctuated strongly between sessions, and most short silences were no more than that. Yield counts from single sessions are therefore noisy, and channels that fall silent for a few sessions are likely to return. Silences lasting over a month were real and still reversed in more than half of cases. Hahn and colleagues analyzed yield, decoding signal quality and tuning stability in the same human data¹⁷. We add electrode-level analyses of silence and recovery against a permutation null, a test of edge-biased spiking loss with the array as the unit, and a link between electrode changes and decoder corrections.

Our study has limitations. All analyses were offline; offline variability can overstate online decay⁴², and in closed loop users adapt to their decoder⁴³. Human decoding labels were inferred from the direction from cursor to target, and the online decoder shaped both the neural activity and the cursor path. Three participants and three monkeys cannot establish population rates, and the datasets differ in task, signal type, array placement and implant age. The label-free comparisons covered three classes of method and did not include latent-dynamics or adversarial alignment. The simulator is calibrated to feature-level statistics in one monkey and does not model biophysics.

Long-term public datasets made this study possible. We release our code, derived results and the calibrated simulator so that others can test methods against realistic drift. Our results suggest that keeping implanted decoders working over months depends less on elaborate corrections than on well-regularized decoders, proper per-channel normalization and short, careful recalibration.

Methods

Datasets and ethics

All data were public and de-identified. The original studies obtained ethical approval and informed consent, as described in the cited publications. We collected no new data.

Monkey N (LINK; DANDI Archive 001201⁴⁴; ref.¹⁹). One rhesus macaque had Utah arrays in the hand area of primary motor cortex³⁵. The release contains 312 sessions on 303 days over 1,242 days, in which the monkey moved two finger groups to targets in center-out or random-target blocks. We used the 96 released channels (64 on the medial array and 32 on the lateral array), with spike-band power and threshold crossings at -4.5 times the root-mean-square noise in 20 ms bins⁴⁵. Outputs were the position and velocity of the two finger groups. Electrode impedance was stored on 228 days.

Monkeys C and M (DANDI Archive 000688⁴⁶; refs.^{28,34}). Both monkeys had Utah arrays in motor and premotor cortex and performed center-out and random-target reaching. Monkey C had 68 sessions from October 2013 to October 2016, and monkey M had 28 sessions from January 2014 to June 2015. We summed the spikes of all sorted units on each electrode in 20 ms bins, giving 96 (C) or 192 (M) channels. Hand position and velocity were interpolated to bin centers.

BrainGate participants (Dryad⁴⁷; ref.¹⁷). The yield data contain one recording from each of 2,289 research sessions in 14 participants with 20 arrays (A1, S1–S3, T1–T3 and T5–T11). For every electrode, they list threshold-crossing rates at several thresholds (with the threshold in microvolts), impedance and grid position. Nineteen arrays were in the hand area of motor cortex and one was in the middle frontal gyrus. The decoding data contain closed-loop cursor sessions with spike-band power in 10 ms bins, cursor and target positions and the periods of attempted movement. We used all sessions of T6 (124 sessions over 3.1 years, one array, 96 channels), T5 (92 sessions over 7.2 years, two arrays, 192 channels) and T9 (117 sessions over 2.0 years, two arrays, 192 channels). Sessions contained a median of 367, 551 and 275 trials, with median movement periods of 1.6, 1.7 and 3.9 s. All sessions of a participant were treated as one task type.

MINDFUL (Dryad⁴⁸; ref.³²). Closed-loop cursor blocks in which T11 (15 days) and T5 (6 days) used a fixed decoder, with neural features in 20 ms bins (spike-band power for T11 and threshold crossings for T5), the decoder’s velocity output and the angular error of the cursor.

Features, decoders and accuracy

For monkey N and the human participants we used spike-band power (human 10 ms bins averaged in pairs). For monkeys C and M we used per-electrode spike counts. Each channel was z-scored with the mean and standard deviation of the decoding session’s training segment, computed without labels (renormalization). The linear decoder was ridge regression⁴⁹ on 8 causal lags (160 ms) of the z-scored features, with an unpenalized intercept. The default penalty was $\alpha = 0.1$ for monkey N and the humans, adopted from ref.¹⁹, and $\alpha = 1$ for monkeys C and M; the tuned penalty was $\alpha = 10^4$ for the decoder trained on day i and for the reference decoder. Outputs were finger position and velocity for monkey N, hand position and velocity for monkeys C and M, and for the humans the intended movement direction, defined as the unit vector from cursor to target during attempted movement. Accuracy was the coefficient of determination R^2 , averaged over outputs and computed on held-out trials of the test session. In monkey N the first 300 of 375 trials trained the decoder, as in ref.¹⁹; in the other datasets the first 80% of trials did.

The recurrent decoder was a long short-term memory network⁵⁰ with 64 hidden units, input dropout⁵¹ of 0.5 and 20-bin input windows (400 ms). It was trained with AdamW⁵² (learning rate 10^{-3} , weight decay 0.01) for up to 4,000 iterations of 256 samples, with early stopping on the last 20% of training trials checked every 100 iterations. Without these measures, a 256-unit network memorized the human training sessions (training R^2 0.97, test R^2 -1.1 on one session). The network analyses used 25 training sessions per individual chosen at random with a fixed seed.

Session pairs

For each training session i and target gap $g \in \{1, 7, 30, 120, 480\}$ days, we selected the later session j of the same task type whose gap was closest to g , within 25% of g or one day, whichever was larger. The one-day set therefore contains one- and two-day gaps, which we report separately. In monkey N, 80 training sessions were drawn at random with a fixed seed; the full monkey N set (537 pairs) also includes further gaps and was used for monitoring. Pair counts are listed in Supplementary Table 2. The time to half retention was found by interpolating median retention at the target gaps linearly in log days, with 95% confidence intervals from a cluster bootstrap over training sessions.

Correction ladder

Let $\mathbf{x}_j(t) \in \mathbb{R}^C$ be the raw features of test day j , $\boldsymbol{\mu}_d$ and $\boldsymbol{\sigma}_d$ the channel means and standard deviations of day d , and f_i the decoder trained on day i .

- *No correction*: $f_i((\mathbf{x}_j - \boldsymbol{\mu}_i)/\boldsymbol{\sigma}_i)$.
- *Mean updating*: $f_i((\mathbf{x}_j - \boldsymbol{\mu}_j)/\boldsymbol{\sigma}_i)$.
- *Renormalization*: $f_i(\mathbf{z}_j)$ with $\mathbf{z}_j = (\mathbf{x}_j - \boldsymbol{\mu}_j)/\boldsymbol{\sigma}_j$.
- *Re-weighting*: f_i applied to $g_c z_{j,c}$ for each channel c , with one offset per output. The weights g_c (unconstrained in sign unless stated) were shrunk toward 1 and the offsets toward the old intercept; the penalty, scaled by the number of samples, was chosen from $\{10^{-6}, \dots, 10^2\}$ by five-fold cross-validation over contiguous blocks of trials.
- *Ridge to prior*: a new decoder $(\mathbf{W}, \mathbf{b})$ fitted on the lagged features $\mathbf{H}$ of day j by minimizing

$$\|\mathbf{Y} - \mathbf{H}\mathbf{W} - \mathbf{1}\mathbf{b}^\top\|_F^2 + \lambda (\|\mathbf{W} - \mathbf{W}_i\|_F^2 + \|\mathbf{b} - \mathbf{b}_i\|_2^2), \quad (1)$$

where $(\mathbf{W}_i, \mathbf{b}_i)$ is the old decoder and λ was chosen from $\{0.1, 1, \dots, 10^6\}$ by cross-validation.

- *Retrained decoder*: ridge regression on day j 's first n labeled trials, with α chosen from the same grid by cross-validation.
- *Reference*: a decoder trained on all of day j 's training trials.

Supervised rungs used the first n labeled trials of day j ($n = 300$, or all training trials if fewer, unless stated). Retention was the ratio of a rung's R^2 to the reference R^2 on the same test trials, and the share of the loss recovered by re-weighting was $(R_{\text{reweight}}^2 - R_{\text{renorm}}^2)/(R_{\text{ref}}^2 - R_{\text{renorm}}^2)$. For the network decoder, re-weighting was a learned gain and offset per channel in front of the frozen network, and the refit was fine-tuning with an L_2 penalty toward the old weights; penalties were chosen on the last 20% of labeled trials. We also computed a full input remap fitted by L-BFGS⁵³; because it can represent any linear decoder of the same form, we used it only as a calibration target for the simulator.

Channel controls. For 40 training sessions per individual (fixed seed), re-weighting was repeated with weights constrained to be non-negative (bounded least squares

with the cross-validated penalty), and for a copy of the day- i decoder whose channel weights were randomly permuted, with signed and with non-negative weights.

Electrode link. For the same pairs in the human participants, the signed re-weighting weights were correlated (Spearman) across channels with the change in log threshold-crossing rate ($\log(\text{rate} + 0.1)$ at -4.5 times the noise) and in log impedance between the yield recordings of the two days, matched by electrode identity. A null correlation was computed after permuting channels.

Matched-output control. For monkeys C and M, sessions were rebuilt with a two-dimensional unit-vector target, the direction of hand velocity, on bins where hand speed exceeded the session’s 25th percentile, and the ladder was run on the same pairs with the tuned penalty.

Decoder-weight angles. Each channel’s block of ridge weights (8 lags by all outputs) was stacked into $\mathbf{B}_i$ and $\mathbf{B}_j$. We computed the angle between $\text{vec}(\mathbf{B}_i)$ and $\text{vec}(\mathbf{B}_j)$, which allows one global scale, and the angle between $\text{vec}(\mathbf{B}_j)$ and $\text{vec}(\text{diag}(\mathbf{s})\mathbf{B}_i)$, with each channel’s scale s_c fitted by least squares.

Label-free corrections

Procrustes rotation. Let $\mathbf{V}_i, \mathbf{V}_j \in \mathbb{R}^{C \times k}$ hold the top $k = 16$ principal axes of the corrected training features on each day. The orthogonal matrix $\mathbf{R}$ minimizing $\|\mathbf{V}_j\mathbf{R} - \mathbf{V}_i\|_F$ was found by Procrustes analysis³⁶, and inputs were mapped as $(\mathbf{I} - \mathbf{V}_j\mathbf{V}_j^\top)\mathbf{z} + \mathbf{V}_i\mathbf{R}^\top\mathbf{V}_j^\top\mathbf{z}$. We applied it to renormalized features and, separately, to mean-updated features with $\mathbf{V}_j$ estimated from those features.

CORAL. Renormalized day- j features were re-colored to match day i ’s covariance, $\mathbf{z}' = \mathbf{C}_i^{1/2}\mathbf{C}_j^{-1/2}\mathbf{z}$ ³⁷, with covariances shrunk 5% toward a scaled identity.

Factor-analysis stabilizer. Following ref.²⁷, a factor-analysis model with 10 latent factors was fitted to each day’s renormalized training features, and a ridge decoder on 8 lags of the day- i posterior latent means was trained. On day j , a model fitted without labels was aligned to day i by orthogonal Procrustes analysis of the loadings of stable channels, found by iteratively keeping the 60% of channels with the smallest loading residuals; latents were then inferred from stable channels only and passed to the frozen latent decoder. The comparison baseline reused the day- i factor model on day j .

All label-free comparisons were summarized per training session over its pairs at all gaps.

Monitoring analyses

In monkey N, for every pair we computed Gaussian Kullback–Leibler divergences between days i and j of the projections of renormalized features on day i ’s top ten principal axes (neural divergence) and of the decoder’s outputs (output divergence). We report raw Spearman correlations with R^2 , partial correlations given log elapsed days and time-blocked cross-validated linear prediction of R^2 . The pilot statistics (first 46–61 sessions, 283 pairs) are described in Supplementary Note 1; split-half reliability

of the residual after the time trend used alternating pairs of trials and the Spearman–Brown formula. For the MINDFUL data we followed the original definitions³²: non-overlapping 60 s windows, divergences against the first day, and the median angular error over non-excluded trials (windows with fewer than 500 valid bins skipped). We report Pearson and partial correlations given elapsed days across windows, and leave-one-day-out linear prediction of angular error; the reference day was excluded from evaluation.

Regularization and recalibration

Ridge decoders were trained with $\alpha \in \{0.1, 10, 10^3, 10^4, 3 \times 10^4, 10^5, 3 \times 10^5, 10^6\}$ on 40 training sessions per individual (22 for monkey M) and evaluated on held-out same-day trials and on later sessions at the five gaps. The 1% rule chose the largest α whose median same-day R^2 was at least 99% of the best median. Compared with each training session’s own best same-day penalty, the rule did not differ in monkeys C and M or in T6 and added 0.004 in monkey N ($P = 6.9 \times 10^{-5}$), 0.000 in T5 and 0.001 in T9. For the human data-efficiency analysis, the ladder was run for 30 training sessions per participant (fixed seed) with $n \in \{10, 20, 50, 100, 300\}$. For harmful recalibrations we used 86 monkey N pairs with default-penalty decoders; the shrinkage λ in equation (1) was chosen per pair by cross-validation on the n trials, or as the value with the best median R^2 over pairs from other training sessions, which was 10^4 for every pair with up to 50 trials. A recalibration was harmful when its R^2 was more than 0.01 below that of the renormalized fixed decoder.

Electrode statistics

Statistics were computed per array. An electrode was active in a session when its threshold-crossing rate at -4.5 times the root-mean-square noise was at least 2 Hz. An initially active electrode (mean rate of at least 2 Hz over the first three sessions) was silenced at its first run of at least three consecutive observed sessions below threshold, and recovered when it later exceeded 4 Hz for at least three consecutive sessions. Missing sessions were skipped, not imputed. Control analyses (Supplementary Table 11) varied these definitions. They used thresholds of -3.5 and -5.5 times the noise, and a fixed threshold per electrode in microvolts: its -4.5 value over the first three sessions, applied by choosing in each session the noise multiple whose microvolt threshold was closest. They also removed sessions with array-wide shifts (absolute median change in log rate above 1), required a mean waveform trough of at least $30 \mu\text{V}$, required silence below 0.5 Hz, or required silences lasting at least 30 days. The null distribution permuted each electrode’s session order independently, 200 times; we report the observed recovery fraction, the null mean and 95% range, and a one-sided permutation P value. For impedance, values of 4,000 k Ω or more were excluded and the trend was the Spearman correlation of the array median with day. For the edge effect, edge electrodes were those on the outer ring of the grid (10×10 in humans and 8×8 in monkey N), each initially active electrode’s decline was the slope of $\log(\text{rate} + 0.1)$ against years, and arrays with at least 10 initially active electrodes were tested with the array as the unit. The direction was predicted from mechanical models³⁹, so we

report one-sided tests; two-sided P values are twice as large. The adjusted analysis regressed slopes on edge membership and initial log rate within each array.

Simulator

The simulator transforms the z-scored features $\mathbf{Z} \in \mathbb{R}^{T \times C}$ of a real base session into those of a simulated session Δt days later (Supplementary Note 2). *Mixing*: $\mathbf{Z} \leftarrow \mathbf{Z}(\mathbf{I} + \mathbf{E})^\top$, with $\mathbf{E} = s(\Delta t)[(1 - \beta)\mathbf{G} + \beta\mathbf{L}]$, where $\mathbf{G}$ is a Gaussian matrix with zero diagonal, $\mathbf{L}$ a Gaussian matrix restricted to grid neighbors, $\beta = 0.5$ and $s(\Delta t)^2 = s_0^2 + [s_\infty(1 - e^{-\Delta t/\tau_s})]^2$. *Turnover*: each channel replaces a fraction ρ_c of its signal by a unit with rotated tuning in the session’s 20-dimensional latent space; ρ_c follows a beta distribution with mean $\rho_0 + (\rho_\infty - \rho_0)(1 - e^{-\Delta t/\tau_\rho})$ and concentration 5. *Silencing*: each channel active in the base session becomes silent with probability $h_{\text{off}}/(h_{\text{off}} + h_{\text{on}})[1 - e^{-(h_{\text{off}} + h_{\text{on}})\Delta t}]$ and keeps 37% of its signal variance, the ratio of tuning strength of silent to active channels in monkey N. The switching rates ($h_{\text{off}} = 0.0038$ and $h_{\text{on}} = 0.0018$ per day) were measured from monkey N’s activity in the calibration period. The six mixing and turnover parameters were calibrated on 12 base sessions before day 700 by minimizing the summed squared difference between simulated and real median retention for three rungs (renormalized, re-weighted and full input remap) at five gaps, with the same correction code and regularization as for real data, using a 24-point Sobol search⁵⁴ followed by Nelder–Mead⁵⁵. Validation used 10 base sessions after day 700 and compared predicted and real medians for all rungs and for $n \in \{20, 50, 100\}$; rungs not fitted were the label-free alignments, the subspace rotation fitted with labels and the two refits. The empirical baseline predicted each late-period median by the corresponding median over pairs whose training session preceded day 700. Augmentation added eight simulated futures per training session (Δt log-uniform between 1 and 120 days) and was compared with stronger regularization, random channel dropout with gain noise²³ and training on the three previous sessions.

Statistics

Session pairs that share a training session are not independent. Confidence intervals for medians therefore came from a cluster bootstrap over training sessions⁵⁶ (2,000 resamples, percentile intervals), and paired tests used one value per training session (two-sided Wilcoxon signed-rank tests). Pairs also share test sessions, so these intervals and P values may be somewhat optimistic; we emphasize effect sizes. Harmful recalibrations were compared with exact McNemar tests on paired outcomes, Holm-adjusted over the five data budgets⁵⁷. With three participants and three monkeys, comparisons between species are descriptive. Exact P values are reported. Analyses used Python 3.11 with NumPy, SciPy, pandas, scikit-learn and PyTorch.

Use of AI tools

An AI assistant (Claude, Anthropic) was used to write analysis code, run analyses and draft the text under the authors’ direction. The authors reviewed the code, analyses and text and take full responsibility for the content.

528 Data availability

529 All data analyzed in this study are public. LINK (monkey N) is available from
530 the DANDI Archive at <https://doi.org/10.48324/dandi.001201/0.251023.2336>. The
531 recordings of monkeys C and M are available at <https://doi.org/10.48324/dandi.000688/0.250122.1735>. The BrainGate yield and decoding data are available from
532 Dryad at <https://doi.org/10.5061/dryad.x0k6djj1h>. The MINDFUL data are avail-
533 able from Dryad at <https://doi.org/10.5061/dryad.n2z34tn5s>. All derived results
534 underlying the figures and tables are provided in the code repository.
535

536 Code availability

537 All code, including the simulator with its calibrated parameters, is available at <https://github.com/Imtiaj-Sajin/invasive-bci> under the MIT license. A versioned archive
538 will be deposited on Zenodo upon publication.
539

540 **Acknowledgements.** We thank the participants of the BrainGate clinical trials
541 and the teams who recorded and released the LINK, DANDI 000688, BrainGate and
542 MINDFUL datasets. This work received no specific funding.

543 **Author contributions.** All authors contributed to the study design, analysis,
544 interpretation and writing, and approved the final manuscript.

545 **Competing interests.** The authors declare no competing interests.

546 References

- 547 [1] Hochberg, L. R. et al. Neuronal ensemble control of prosthetic devices by a human
548 with tetraplegia. *Nature* **442**, 164–171 (2006).
- 549 [2] Hochberg, L. R. et al. Reach and grasp by people with tetraplegia using a neurally
550 controlled robotic arm. *Nature* **485**, 372–375 (2012).
- 551 [3] Collinger, J. L. et al. High-performance neuroprosthetic control by an individual
552 with tetraplegia. *Lancet* **381**, 557–564 (2013).
- 553 [4] Pandarinath, C. et al. High performance communication by people with paralysis
554 using an intracortical brain–computer interface. *eLife* **6**, e18554 (2017).
- 555 [5] Willett, F. R., Avansino, D. T., Hochberg, L. R., Henderson, J. M. & Shenoy, K.
556 V. High-performance brain-to-text communication via handwriting. *Nature* **593**,
557 249–254 (2021).
- 558 [6] Willett, F. R. et al. A high-performance speech neuroprosthesis. *Nature* **620**,
559 1031–1036 (2023).
- 560 [7] Card, N. S. et al. An accurate and rapidly calibrating speech neuroprosthesis. *N.*
561 *Engl. J. Med.* **391**, 609–618 (2024).

- 562 [8] Wairagkar, M. et al. An instantaneous voice-synthesis neuroprosthesis. *Nature*
563 **644**, 145–152 (2025).
- 564 [9] Card, N. S. et al. Long-term independent use of an intracortical brain–computer
565 interface for speech and cursor control. *Nat. Med.* **32**, 2504–2510 (2026).
- 566 [10] Perge, J. A. et al. Intra-day signal instabilities affect decoding performance in an
567 intracortical neural interface system. *J. Neural Eng.* **10**, 036004 (2013).
- 568 [11] Chestek, C. A. et al. Long-term stability of neural prosthetic control signals from
569 silicon cortical arrays in rhesus macaque motor cortex. *J. Neural Eng.* **8**, 045005
570 (2011).
- 571 [12] Downey, J. E., Schwed, N., Chase, S. M., Schwartz, A. B. & Collinger, J. L. Intra-
572 cortical recording stability in human brain–computer interface users. *J. Neural*
573 *Eng.* **15**, 046016 (2018).
- 574 [13] Bishop, W. E. et al. Self-recalibrating classifiers for intracortical brain–computer
575 interfaces. *J. Neural Eng.* **11**, 026001 (2014).
- 576 [14] Barrese, J. C. et al. Failure mode analysis of silicon-based intracortical micro-
577 electrode arrays in non-human primates. *J. Neural Eng.* **10**, 066014 (2013).
- 578 [15] Sponheim, C. et al. Longevity and reliability of chronic unit recordings using the
579 Utah, intracortical multi-electrode arrays. *J. Neural Eng.* **18**, 066044 (2021).
- 580 [16] Bjånes, D. A. et al. Quantifying physical degradation alongside recording
581 and stimulation performance of 980 intracortical microelectrodes chronically
582 implanted in three humans for 956–2130 days. *Acta Biomater.* **198**, 188–206
583 (2025).
- 584 [17] Hahn, N. V. et al. Performance of intracortical microelectrode arrays in people
585 with implanted brain–computer interfaces over a 20-year period. *Nat. Med.* (in
586 the press); preprint at <https://doi.org/10.1101/2025.07.02.25330310> (2025).
- 587 [18] Nuyujukian, P. et al. Performance sustaining intracortical neural prostheses. *J.*
588 *Neural Eng.* **11**, 066003 (2014).
- 589 [19] Temmar, H. et al. Long-term Intracortical Neural activity and Kinematics
590 (LINK): an intracortical neural dataset for chronic brain–machine interfaces,
591 neuroscience, and machine learning. *Adv. Neural Inf. Process. Syst.* **38**, 154225–
592 154247 (2025).
- 593 [20] Wilson, G. H. et al. Long-term unsupervised recalibration of cursor-based intra-
594 cortical brain–computer interfaces using a hidden Markov model. *Nat. Biomed.*
595 *Eng.* **10**, 1466–1484 (2025).

- 596 [21] Jarosiewicz, B. et al. Virtual typing by people with tetraplegia using a self-
597 calibrating intracortical brain–computer interface. *Sci. Transl. Med.* **7**, 313ra179
598 (2015).
- 599 [22] Brandman, D. M. et al. Rapid calibration of an intracortical brain–computer
600 interface for people with tetraplegia. *J. Neural Eng.* **15**, 026007 (2018).
- 601 [23] Sussillo, D., Stavisky, S. D., Kao, J. C., Ryu, S. I. & Shenoy, K. V. Making brain–
602 machine interfaces robust to future neural variability. *Nat. Commun.* **7**, 13749
603 (2016).
- 604 [24] Orsborn, A. L., Dangi, S., Moorman, H. G. & Carmena, J. M. Closed-loop
605 decoder adaptation on intermediate time-scales facilitates rapid BMI performance
606 improvements independent of decoder initialization conditions. *IEEE Trans.*
607 *Neural Syst. Rehabil. Eng.* **20**, 468–477 (2012).
- 608 [25] Dangi, S., Orsborn, A. L., Moorman, H. G. & Carmena, J. M. Design and analysis
609 of closed-loop decoder adaptation algorithms for brain–machine interfaces. *Neural*
610 *Comput.* **25**, 1693–1731 (2013).
- 611 [26] Fan, C. et al. Plug-and-play stability for intracortical brain–computer interfaces:
612 a one-year demonstration of seamless brain-to-text communication. *Adv. Neural*
613 *Inf. Process. Syst.* **36** (2023).
- 614 [27] Degenhart, A. D. et al. Stabilization of a brain–computer interface via the align-
615 ment of low-dimensional spaces of neural activity. *Nat. Biomed. Eng.* **4**, 672–685
616 (2020).
- 617 [28] Gallego, J. Á., Perich, M. G., Chowdhury, R. H., Solla, S. A. & Miller, L. E. Long-
618 term stability of cortical population dynamics underlying consistent behavior.
619 *Nat. Neurosci.* **23**, 260–270 (2020).
- 620 [29] Farshchian, A. et al. Adversarial domain adaptation for stable brain–machine
621 interfaces. In *Proc. 7th International Conference on Learning Representations*
622 (2019).
- 623 [30] Ma, X. et al. Using adversarial networks to extend brain computer interface
624 decoding accuracy over time. *eLife* **12**, e84296 (2023).
- 625 [31] Karpowicz, B. M. et al. Stabilizing brain–computer interfaces through alignment
626 of latent dynamics. *Nat. Commun.* **16**, 4662 (2025).
- 627 [32] Pun, T. K. et al. Measuring instability in chronic human intracortical neural
628 recordings towards stable, long-term brain–computer interfaces. *Commun. Biol.*
629 **7**, 1363 (2024).

- [33] Karpowicz, B. M. et al. Few-shot algorithms for consistent neural decoding (FALCON) benchmark. *Adv. Neural Inf. Process. Syst.* **37** (2024).
- [34] Perich, M. G., Gallego, J. Á. & Miller, L. E. A neural population mechanism for rapid learning. *Neuron* **100**, 964–976.e7 (2018).
- [35] Nason, S. R. et al. Real-time linear prediction of simultaneous and independent movements of two finger groups using an intracortical brain–machine interface. *Neuron* **109**, 3164–3177.e8 (2021).
- [36] Schönemann, P. H. A generalized solution of the orthogonal Procrustes problem. *Psychometrika* **31**, 1–10 (1966).
- [37] Sun, B., Feng, J. & Saenko, K. Return of frustratingly easy domain adaptation. In *Proc. 30th AAAI Conference on Artificial Intelligence* 2058–2065 (AAAI Press, 2016).
- [38] Chen, X. et al. Chronic stability of a neuroprosthesis comprising multiple adjacent Utah arrays in monkeys. *J. Neural Eng.* **20**, 036039 (2023).
- [39] Forrest, A. M. et al. Finite element model predicts micromotion-induced strain profiles that correlate with the functional performance of Utah arrays in humans and non-human primates. *J. Neural Eng.* **22**, 066008 (2025).
- [40] Polikov, V. S., Tresco, P. A. & Reichert, W. M. Response of brain tissue to chronically implanted neural electrodes. *J. Neurosci. Methods* **148**, 1–18 (2005).
- [41] Kozai, T. D. Y., Jaquins-Gerstl, A., Vazquez, A. L., Michael, A. C. & Cui, X. T. Brain tissue responses to neural implants impact signal sensitivity and intervention strategies. *ACS Chem. Neurosci.* **6**, 48–67 (2015).
- [42] Flint, R. D., Wright, Z. A., Scheid, M. R. & Slutzky, M. W. Long term, stable brain machine interface performance using local field potentials and multiunit spikes. *J. Neural Eng.* **10**, 056005 (2013).
- [43] Gilja, V. et al. A high-performance neural prosthesis enabled by control algorithm design. *Nat. Neurosci.* **15**, 1752–1757 (2012).
- [44] Temmar, H. et al. LINK: Long-Term Intracortical Neural Activity and Kinematics. *DANDI Archive* <https://doi.org/10.48324/dandi.001201/0.251023.2336> (2025).
- [45] Nason, S. R. et al. A low-power band of neuronal spiking activity dominated by local single units improves the performance of brain–machine interfaces. *Nat. Biomed. Eng.* **4**, 973–983 (2020).
- [46] Perich, M. G., Miller, L. E., Azabou, M. & Dyer, E. L. Long-term recordings of motor and premotor cortical spiking activity during reaching in monkeys. *DANDI*

- 665 *Archive* <https://doi.org/10.48324/dandi.000688/0.250122.1735> (2025).
- 666 [47] Hahn, N. V. et al. Data from: Performance of intracortical microelectrode arrays
667 in people with implanted brain–computer interfaces over a 20-year period. *Dryad*
668 <https://doi.org/10.5061/dryad.x0k6djj1h> (2026).
- 669 [48] Pun, T. K. et al. Data from: Measuring instability in chronic human intracortical
670 neural recordings towards stable, long-term brain–computer interfaces. *Dryad*
671 <https://doi.org/10.5061/dryad.n2z34tn5s> (2024).
- 672 [49] Hoerl, A. E. & Kennard, R. W. Ridge regression: biased estimation for nonorthog-
673 onal problems. *Technometrics* **12**, 55–67 (1970).
- 674 [50] Hochreiter, S. & Schmidhuber, J. Long short-term memory. *Neural Comput.* **9**,
675 1735–1780 (1997).
- 676 [51] Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I. & Salakhutdinov, R.
677 Dropout: a simple way to prevent neural networks from overfitting. *J. Mach.*
678 *Learn. Res.* **15**, 1929–1958 (2014).
- 679 [52] Loshchilov, I. & Hutter, F. Decoupled weight decay regularization. In *Proc. 7th*
680 *International Conference on Learning Representations* (2019).
- 681 [53] Liu, D. C. & Nocedal, J. On the limited memory BFGS method for large scale
682 optimization. *Math. Program.* **45**, 503–528 (1989).
- 683 [54] Sobol’, I. M. On the distribution of points in a cube and the approximate
684 evaluation of integrals. *USSR Comput. Math. Math. Phys.* **7**, 86–112 (1967).
- 685 [55] Nelder, J. A. & Mead, R. A simplex method for function minimization. *Comput.*
686 *J.* **7**, 308–313 (1965).
- 687 [56] Field, C. A. & Welsh, A. H. Bootstrapping clustered data. *J. R. Stat. Soc. B* **69**,
688 369–390 (2007).
- 689 [57] Holm, S. A simple sequentially rejective multiple test procedure. *Scand. J. Stat.*
690 **6**, 65–70 (1979).