Supplementary Information for: Decoder decay and its correction over years in intracortical brain–computer interfaces

Md. Imtiaj Alam Sajin[1*], Esm E Moula Chowdhury Abha[1] and Md Wahiduzzaman Suva[1]

[1*]Department of Computer Science, American International University-Bangladesh, Dhaka, 1229, Bangladesh.

*Corresponding author(s). E-mail(s): imtiajsajin@gmail.com;
Contributing authors: esmechowdhuryabha@gmail.com;
wshuvo360@gmail.com;

Abstract

This document contains Supplementary Notes 1–5 and Supplementary Tables 1–18. Tables 2–18 are generated from the released result files by `scripts/make_supp_tables.py`, so their values match the code repository; Table 1 lists the datasets as described by their sources.

Supplementary Note 1: Label-free monitoring statistics

Before the main analyses, we asked whether label-free statistics computed on a later day can predict how well a renormalized fixed decoder will perform on that day. The pilot used the first 46–61 sessions of monkey N (January 2020 to about March 2021), with session pairs up to more than 180 days apart; the cross-validated prediction below used the 283 pairs at most 60 days apart. Decoders were ridge regressions on 8 lags of spike-band power ($\alpha = 0.1$), trained on the first 300 trials and tested on the rest.

We computed 17 statistics in nine groups that need no labels on the later day:

- the shift in channel means and in channel standard deviations;
- the principal angles between the dominant subspaces of the two days;

- disagreement within a bootstrap ensemble of decoders;
- disagreement between decoders built on different views of the data (spike-band power versus threshold crossings, and 8 lags versus 1 lag);
- statistics of the predicted output (velocity bias, variance ratio, Wasserstein distance to the training kinematics and smoothness);
- self-consistency between predicted position and predicted velocity;
- Procrustes alignment disagreement, the Procrustes residual and the share of decoder readout outside the dominant subspace;
- neural modulation depth;
- the log ratio of impedances and the number of silent channels.

Elapsed time predicted decoder accuracy well (Spearman $\rho = -0.75$ across all pairs and -0.56 within 60 days). After controlling for log elapsed days, the partial Spearman correlation of every statistic with R^2 was at most 0.27 (output variance) and mostly below 0.1. In time-blocked cross-validated prediction of R^2, elapsed days alone gave a mean absolute error of 0.064. The label-free statistics alone gave 0.067–0.071, and elapsed days together with the statistics gave 0.067–0.070. When the target was the gain from recalibration rather than R^2, elapsed days gave 0.045 and the statistics gave 0.054. For detecting unexpectedly poor days (residual more than one standard deviation below the time trend), every statistic had an area under the receiver operating characteristic curve of 0.50–0.58.

None of the label-free statistics we tried captured day-to-day performance beyond the time trend. Abrupt faults on single channels remain detectable by construction, and such faults may still justify monitoring.

Supplementary Note 2: Simulator

The simulator (`ibci.sim`) is a feature-level generative model of chronic change. It transforms a real base session into a simulated session Δt days later, so that behavior and latent dynamics stay realistic. Its components are defined in the Methods and its parameters are listed in Supplementary Table 12. Two further components, heavy-tailed gain jumps and random walks of raw offsets and scales, are included for decoders that are not renormalized. They do not affect renormalized decoders and were not calibrated.

Calibration. The simulated ladder must be scored with exactly the same correction code and regularization as the real ladder. An earlier calibration used regularization 100–1,000 times stronger than the values chosen on real data. That handicapped the simulated corrections, and the optimizer compensated with distorted parameters (loss 0.663 versus 0.074 after the fix). We also restricted silencing to channels that were active in the base session, because most channels are already silent on any given day.

Validation. The simulator calibrated on days before 700 was evaluated on base sessions after day 700, for all rungs of the ladder and for 20, 50 and 100 labeled trials. Two references were used: the simulator with all drift switched off, and the median retention measured on real pairs before day 700 (an empirical curve). The simulator beat the no-drift reference everywhere. It beat the empirical curve only for corrections it was not fitted to; for fitted corrections and for different amounts of labeled data the

empirical curve was as good or better (main text, Fig. 8b). The simulator is therefore most useful for predicting how untested corrections behave.

Intended uses and limits. The simulator is meant for stress-testing decoders and label-free methods under controlled drift and electrode switching, and for comparing recalibration policies. It is calibrated on one monkey with Utah arrays, spike-band power and a finger task, and it has not been calibrated to human recordings. Its mechanisms reproduce the measured statistics, which does not show that they are the true physical mechanisms. Users can calibrate it to other datasets with `scripts/calibrate_sim.py`.

Supplementary Note 3: Analysis choices revised during the study

We report analysis choices that we revised after internal review, because each could mislead similar studies.

One-day pairs included two-day gaps. Session pairs were matched to a target gap within 25% or one day, whichever was larger, so the one-day set contains two-day gaps. In participant T6 most "one-day" pairs were two-day pairs, which changed the apparent overnight retention. We now report one- and two-day pairs separately.

Decoder regularization confounds decay. With the small penalties we first adopted ($\alpha = 0.1$ from the monkey N study, $\alpha = 1$ for monkeys C and M), decoders decayed faster than with a tuned penalty, most in monkey N, T5 and T7. Decay is therefore reported for decoders whose penalty was tuned.

Label-free alignment was compared with renormalization. In our ladder, alignment was first applied to renormalized features, which cannot separate the two. We therefore also aligned mean-updated features.

Electrode revival needs a null. Counting electrodes that fall silent and later recover suggested a striking recovery rate. The same statistic computed on shuffled session orders was almost as high, so only long silences carry evidence of a temporal process. Electrode-level tests also treat the array, not the electrode, as the unit, because neighboring electrodes are correlated.

Flexible corrections need scrambled-channel controls. Re-weighting one weight per channel recovered most of the loss in the human participants, which at first suggested that drift there is mostly per-channel. Decoders with scrambled channels could also be re-weighted to high accuracy when the decoded variable was two-dimensional, so high recovery alone does not show per-channel drift. We therefore report scrambled-channel and non-negative controls (Supplementary Table 5).

A full input remap is not a structural test. A free $C \times C$ linear map in front of a frozen decoder can reproduce any decoder of the same form when the readout rank is below the number of channels. We used it only as a calibration target for the simulator.

Recurrent networks memorize closed-loop sessions. A 256-unit network trained without early stopping reached training R^2 of 0.97 and test R^2 of -1.1 on a human session. All network results use a smaller, regularized network with early stopping on held-out training trials.

Small-budget recalibration shrinks the intercept. Refitting the intercept freely with 10–20 trials shifted the output offset and made recalibration harmful in many pairs.

Supplementary Note 4: Excluded participants and sensitivity analyses

Participants whose decoders did not work. The BrainGate decoding data contain nine participants. Retention is defined only when the reference decoder, trained on the test day, works. We therefore required a median reference R^2 above 0.1 across a participant's session pairs (Methods). The rule was fixed when the first added participant (T3) failed and before the results of the other added participants were seen. T3 (8 sessions, 12 pairs) and T2 (42 sessions, 120 pairs) failed it, with median reference R^2 of -0.05 and -0.02. Their movement periods lasted a median of 15.6 s, against 1.6–3.9 s in the included participants (Supplementary Table 14). For T2, block-wise normalization and threshold crossings did not raise same-day R^2 above about 0.1 except on its earliest session, whereas the same checks left T6 and T7 working.

T11. T11 (196 sessions, 677 pairs) also failed the rule (median reference R^2 -16.19; 155 pairs above 0.1). In 107 of 196 sessions, the held-out trials contained spike-band power more than 10^3 standard deviations from the training trials of the same session. This was rare in the included participants (T7 2 of 35, T8 4 of 58, T10 1 of 57 sessions), and such pairs are removed by the pair rule below. The jumps coincided with recording blocks. When each block was z-scored separately, as in the dataset paper (Hahn et al., 2026), most affected T11 sessions decoded normally, and clean sessions were unchanged. Block-wise z-scoring uses the statistics of every block, including blocks in the evaluation segment, so this analysis uses more unlabeled information from the test day than renormalization does. With block-wise normalization, the median reference R^2 of T11 was 0.14 (351 of 677 pairs above 0.1). The renormalized fixed decoder kept a median retention of 0.53 after one day, -0.06 after 30 days and -0.49 after 480 days, and re-weighting each channel restored retention to 0.86–0.99, in line with the included participants. T11 is not part of the main analysis because the inclusion rule was fixed before its results were seen.

Session pairs with a failed reference decoder. Within included individuals, pairs whose reference R^2 was at most 0.1 were left out (0–18% of pairs per individual). With all pairs kept, the main conclusions did not change. Time to half retention ranged from <1 (T8) to 125 days (T6), compared with <1 to 138 days with the rule, so the spread between individuals remained wide. The largest relative change was in T7, whose retention at the first gap (one or two days) lay close to one half (0.59 with all pairs, 0.48 with the rule), so its time to half retention moved from 7.1 days to under one day. T6 changed from 125 to 138 days, which is small relative to its slow decay, and the other individuals changed by less than one day. Full results with all pairs are in `results/v2_summary_allpairs.json`.

Do decay curves differ between individuals? Half-lives are censored below one day and have wide intervals, so we also fitted a linear mixed model (Methods; `scripts/heterogeneity_model.py`). Whether the curves differ was clear under every outcome definition. With retention floored at 0, the likelihood-ratio statistic was $\chi^2_{16} =$

174.6 ($P = 1.0 \times 10^{-28}$) for all nine individuals and $\chi^2_{10} = 116.3$ ($P = 2.8 \times 10^{-20}$) for the participants. With a floor of -1, it was 131.3 ($P = 3.6 \times 10^{-20}$) and 83.4 ($P = 1.1 \times 10^{-13}$). For the absolute loss in R^2, it was 61.1 ($P = 3.4 \times 10^{-7}$) and 191.2 ($P = 1.1 \times 10^{-35}$). Whether the slopes also differ, beyond the level of each curve, was less clear. For all nine individuals the slope test gave $P = 2.1 \times 10^{-6}$ (floor 0), 0.002 (floor -1) and 0.008 (absolute loss). For the participants alone it gave $P = 0.059$, 0.004 and 1.4×10^{-6}. The floor at 0 compresses the slopes of participants whose decoders often fail, which may explain the weaker result with that outcome. We therefore claim only that the curves differ. The mixed model assumes one residual variance for all individuals, which their different scales may violate. As a check without that assumption, we took one value per training session (the mean retention of its pairs) at each gap and compared individuals with a Kruskal–Wallis test. Individuals differed at every gap of 7 days or more ($P \leq 7.2 \times 10^{-7}$; participants alone $P \leq 2.1 \times 10^{-5}$). After one day the difference was smaller ($P = 0.005$), and among the participants it was not detectable ($P = 0.19$).

Supplementary Note 5: Pretrained foundation decoder

Pretraining data. The public NDT3 checkpoint pretrained on about 2,000 hours (45 million parameters) stores the list of its pretraining sessions: 26,298 sessions from 23 subjects, mostly human sessions from the University of Pittsburgh plus monkey reaching datasets. The monkey datasets include those of monkeys C and M of this study, so we did not test them. No BrainGate session and no session of monkey N is listed.

Choice of features and individuals. NDT3 takes threshold-crossing counts. With threshold crossings, the same-day ridge decoder reached a median R^2 of 0.22 in monkey N (0.29 with spike-band power), 0.50 in T5 (0.54) and 0.32 in T8 (0.34). The crossings of T6 at -4.5 times the RMS were too sparse to decode (0.04, against 0.28 with spike-band power), so T6 was not used. We chose monkey N and the two participants with the best threshold-crossing decoding and the longest recordings among those whose decoders worked well.

Learning rate. A first run used a learning rate of 10^{-4}. In T5, NDT3's same-day R^2 then fell below the ridge decoder in most sessions. On one T5 session (ridge decoder R^2 0.49), a learning rate of 4×10^{-4} gave 0.38, against 0.26 with 10^{-4}; longer training gave 0.30, and a shorter context during evaluation 0.29. Repeated runs with identical settings differed by up to 0.07. We then fixed the rule described in Methods: the rate with the higher NDT3 validation R^2 on each individual's earliest session (monkey N 0.44 with 4×10^{-4} versus 0.42; T5 0.61 versus 0.59; T8, on the array kept by the array rule, 0.37 with 10^{-4} versus 0.34). These differences are smaller than the run-to-run variation, so the chosen rate is itself uncertain; this is one reason we report NDT3 as a pilot. Monkey N and T5 were rerun completely with the chosen rate. In monkey N, the first run (10^{-4}) gave the same pattern: same-day R^2 0.39 and median retention 1.08 with 300 trials.

A nearly silent array in T8. With both arrays, NDT3 did not learn T8 (R^2 0.00 on its earliest session). Training the same architecture from scratch gave -0.01, and

the first array alone −0.02. With the second array alone, NDT3 reached 0.31, against 0.38 for the ridge decoder on all channels. Only 17 of the first array's 96 channels were active, and the ridge decoder reached 0.02 on that array alone and 0.38 with or without it, whereas NDT3 failed whenever it was present. We therefore fixed the array rule in Methods before running T8's pairs. It removed this array only.

Evaluation. Predictions were made in overlapping 2-s windows advancing by 1 s, keeping the last second of each. On a monkey N session this gave R^2 of 0.39 against 0.40 for NDT3's streaming evaluation with 1 s of context, and 0.08 against 0.09 seven days later. NDT3's data loader drops windows whose behaviour is constant, which happens when the cursor-to-target direction does not change for 2 s. Because behaviour inputs are masked during prediction, we added a ramp of 10^{-3} to the behaviour written for evaluation windows so that every window was kept; predictions were unchanged.

Scope. The pilot used one model size, threshold crossings and 12 session pairs in each of three individuals, all offline. A rule fixed before the pilot required NDT3 to match or exceed the ridge decoder on same-day data in at least two of three individuals before a full benchmark run. It was met only in monkey N, so the analysis was not extended. Larger models, other adaptation schemes or more labeled data may behave differently.

Supplementary Table 1 Datasets. Session counts are per analysis; some human sessions contribute to more than one analysis. SBP, spike-band power; TC, threshold crossings. [a]Recorded on 303 days.

Individual	Source	Arrays and site	Signal	Task	Sessions (span)
Monkey N	DANDI 001201 (LINK)	2 Utah, motor cortex	SBP, TC	finger flexion	312[a] (3.4 years)
Monkey C	DANDI 000688	Utah, motor	sorted spikes	reaching	68 (3.0 years)
Monkey M	DANDI 000688	Utah, motor/premotor	sorted spikes	reaching	28 (1.5 years)
Human T6	Dryad x0k6djj1h	1 Utah, motor cortex	SBP	closed-loop cursor	124 (3.1 years)
Human T5	Dryad x0k6djj1h	2 Utah, motor cortex	SBP	closed-loop cursor	92 (7.2 years)
Human T9	Dryad x0k6djj1h	2 Utah, motor cortex	SBP	closed-loop cursor	117 (2.0 years)
Human T7	Dryad x0k6djj1h	2 Utah, motor cortex	SBP	closed-loop cursor	35 (1.4 years)
Human T10	Dryad x0k6djj1h	1 Utah, motor cortex	SBP	closed-loop cursor	57 (0.9 years)
Human T8	Dryad x0k6djj1h	2 Utah, motor cortex	SBP	closed-loop cursor	58 (2.7 years)
Human T2 (excluded)	Dryad x0k6djj1h	1 Utah, motor cortex	SBP	closed-loop cursor	42 (2.0 years)
Human T3 (excluded)	Dryad x0k6djj1h	1 Utah, motor cortex	SBP	closed-loop cursor	8 (1.0 years)
Human T11 (excluded)	Dryad x0k6djj1h	2 Utah, motor cortex	SBP	closed-loop cursor	196 (4.6 years)
14 humans	Dryad x0k6djj1h	20 Utah arrays	TC rates, impedance	research sessions	2,289 ($\leq$7.3 years)
T11, T5	Dryad n2z34tn5s (MINDFUL)	motor cortex	SBP (T11), TC (T5)	closed-loop cursor	15 and 6 days

Supplementary Table 2 Correction ladder by individual and gap (linear decoder, tuned penalty $\alpha = 10^4$). Median decoding R^2 for each rung, and median retention of the renormalized fixed decoder with 95% cluster-bootstrap confidence intervals. Each training session contributes at most one pair per gap. Supervised rungs used up to 300 labeled trials. Aligned, Procrustes rotation after renormalization; stable aligned, the same estimated from stable channels. Monkeys here; the human participants follow on the next page.

Gap (days)	Pairs	Reference	None	Mean update	Renorm.	Aligned	Stable aligned	Re-weighted	Ridge to prior	Retention [95% CI]
Monkey N										
1	30	0.29	−0.61	0.19	0.24	0.21	0.21	0.23	0.31	0.79 [0.74, 0.85]
7	48	0.30	−0.75	0.16	0.19	0.16	0.16	0.20	0.30	0.63 [0.60, 0.68]
30	61	0.28	−1.46	0.07	0.11	0.11	0.08	0.12	0.29	0.41 [0.34, 0.57]
120	70	0.29	−1.48	0.06	0.08	0.06	0.05	0.09	0.29	0.30 [0.25, 0.35]
480	50	0.28	−0.95	0.03	0.00	−0.01	−0.00	0.08	0.29	0.01 [−0.03, 0.09]
Monkey C										
1	37	0.27	0.10	0.20	0.20	0.19	−0.10	0.24	0.28	0.73 [0.65, 0.80]
7	31	0.33	0.02	0.20	0.22	0.19	−0.06	0.27	0.33	0.62 [0.55, 0.70]
30	11	0.45	−0.15	0.10	0.10	0.11	0.09	0.15	0.45	0.23 [0.19, 0.55]
120	18	0.18	−0.61	−0.00	0.01	−0.03	−0.21	0.07	0.18	0.04 [−0.16, 0.17]
480	33	0.15	−0.78	−0.15	−0.06	−0.29	−0.36	0.02	0.14	−0.12 [−0.42, −0.10]
Monkey M										
1	13	0.36	0.17	0.27	0.27	0.25	−0.02	0.31	0.37	0.80 [0.74, 0.86]
7	6	0.27	0.06	0.15	0.18	0.13	−0.11	0.19	0.27	0.62 [0.51, 0.70]
30	6	0.43	−0.27	0.22	0.26	0.27	0.24	0.30	0.44	0.57 [0.48, 0.68]
120	9	0.25	−0.23	0.00	0.02	0.01	−0.22	0.07	0.22	0.09 [−0.07, 0.19]
480	9	0.35	−0.23	−0.03	0.01	−0.07	−0.15	0.07	0.35	0.04 [−0.47, 0.11]

Supplementary Table 2 (continued) Correction ladder, human participants. Columns as on the previous page.

Gap (days)	Pairs	Reference	None	Mean update	Renorm.	Aligned	Stable aligned	Re-weighted	Ridge to prior	Retention [95% CI]
Human T6										
1	23	0.31	−0.58	0.17	0.18	0.20	0.20	0.27	0.30	0.79 [0.55, 0.84]
7	57	0.25	−0.36	0.10	0.15	0.15	0.14	0.20	0.22	0.74 [0.63, 0.84]
30	79	0.26	−0.52	0.08	0.20	0.21	0.20	0.26	0.28	0.76 [0.70, 0.83]
120	98	0.26	−0.74	0.04	0.16	0.14	0.13	0.23	0.26	0.60 [0.48, 0.66]
480	76	0.30	−1.37	0.03	0.12	0.12	0.12	0.26	0.29	0.35 [0.26, 0.44]
Human T5										
1	17	0.56	−1.02	0.12	0.42	0.47	0.48	0.53	0.51	0.73 [0.34, 0.83]
7	25	0.55	−1.70	−0.37	0.26	0.26	0.28	0.45	0.51	0.52 [0.23, 0.68]
30	49	0.48	−3.88	−0.85	0.13	0.10	0.12	0.39	0.46	0.28 [−0.02, 0.56]
120	72	0.53	−3.73	−1.13	0.07	−0.01	−0.08	0.42	0.50	0.13 [−0.02, 0.21]
480	84	0.52	−6.10	−0.96	−0.09	−0.15	−0.13	0.42	0.49	−0.17 [−0.28, −0.06]
Human T9										
1	23	0.32	−0.92	−0.01	0.18	0.14	0.14	0.27	0.31	0.60 [0.40, 0.68]
7	53	0.33	−0.97	−0.10	0.11	0.10	0.10	0.28	0.31	0.40 [0.32, 0.51]
30	103	0.35	−4.08	−0.37	0.10	0.07	0.10	0.27	0.33	0.31 [0.15, 0.43]
120	90	0.36	−4.46	−0.31	0.04	−0.02	0.00	0.25	0.33	0.14 [0.01, 0.19]
480	51	0.36	−7.40	−0.45	−0.15	−0.15	−0.18	0.25	0.34	−0.38 [−0.53, −0.35]
Human T7										
1	4	0.35	−1.59	−0.88	0.20	0.14	0.11	0.28	0.35	0.53 [0.20, 0.72]
7	8	0.28	−2.69	−0.30	0.15	0.15	0.14	0.23	0.26	0.51 [−0.09, 0.71]
30	19	0.31	−2.77	−1.22	0.09	0.09	0.06	0.25	0.31	0.29 [−0.02, 0.46]
120	30	0.34	−6.88	−1.24	−0.02	−0.05	−0.09	0.26	0.32	−0.08 [−0.54, 0.13]
480	12	0.57	−2.46	−0.74	0.03	−0.10	−0.06	0.44	0.57	0.06 [−0.20, 0.09]
Human T10										
1	16	0.34	−0.06	−0.01	0.20	0.20	0.17	0.24	0.31	0.58 [0.04, 0.79]
7	26	0.30	−1.35	−0.14	0.12	0.11	0.12	0.22	0.28	0.42 [0.27, 0.65]
30	43	0.27	−1.42	−0.37	0.06	0.08	0.04	0.19	0.24	0.18 [0.05, 0.32]
120	39	0.27	−3.34	−0.82	−0.01	−0.01	0.01	0.20	0.23	−0.04 [−0.41, 0.16]
Human T8										
1	7	0.31	−2.16	−0.34	0.13	0.01	0.12	0.17	0.20	0.49 [0.23, 0.82]
7	21	0.34	−5.23	−0.61	0.05	−0.01	−0.01	0.10	0.12	0.19 [−0.12, 0.52]
30	32	0.28	−5.46	−0.60	−0.01	−0.08	−0.06	0.16	0.17	−0.03 [−0.53, 0.30]
120	48	0.30	−6.80	−1.10	−0.00	−0.06	−0.05	0.15	0.21	−0.02 [−0.21, 0.29]
480	47	0.43	−11.79	−1.65	−0.11	−0.18	−0.21	0.29	0.41	−0.24 [−0.49, −0.13]

Supplementary Table 3 Label-free corrections. Median over training sessions of each session's mean paired difference in R^2 (its pairs at all gaps, 1 to 480 days; pairs with a working reference decoder), with two-sided Wilcoxon signed-rank P in parentheses. Mean, mean updating only; rot., Procrustes rotation; stab., factor-analysis stabilizer; FA, factor analysis.

Individual	Renorm. vs mean	Mean+rot. vs mean	Renorm.+rot. vs renorm.	CORAL vs renorm.	Stab. vs renorm.	Stab. vs fixed FA
Monkey N	0.02 (1e-10)	−0.06 (4e-11)	−0.02 (4e-11)	−0.00 (0.95)	−0.06 (9e-14)	−0.02 (2e-11)
Monkey C	0.02 (2e-4)	−0.07 (6e-10)	−0.02 (7e-8)	−0.00 (9e-5)	−0.11 (8e-6)	−0.06 (9e-9)
Monkey M	0.01 (5e-5)	−0.08 (1e-6)	−0.03 (3e-4)	−0.00 (0.77)	−0.15 (5e-7)	−0.09 (5e-7)
Human T6	0.09 (2e-19)	−0.07 (6e-11)	−0.01 (5e-5)	−0.01 (2e-8)	−0.08 (4e-16)	−0.06 (9e-17)
Human T5	2.10 (3e-16)	−9.19 (2e-15)	−0.05 (1e-5)	−0.08 (1e-15)	0.02 (0.038)	0.02 (0.14)
Human T9	0.44 (8e-20)	−2.56 (1e-18)	−0.05 (1e-12)	−0.02 (3e-12)	−0.11 (2e-11)	−0.03 (2e-7)
Human T7	1.34 (5e-10)	−8.69 (5e-10)	−0.03 (0.011)	−0.03 (5e-6)	−0.03 (0.17)	0.02 (0.2)
Human T10	0.88 (6e-10)	−4.40 (5e-10)	−0.00 (0.25)	−0.02 (1e-8)	−0.03 (0.0028)	−0.02 (0.014)
Human T8	1.74 (1e-10)	−2.24 (4e-6)	−0.06 (4e-9)	−0.01 (3e-4)	−0.09 (0.014)	0.00 (0.93)

a

Supplementary Table 4 Monitoring. a, Offline, monkey N (537 pairs): Spearman correlation with decoder R^2, raw and partial given log elapsed days, and time-blocked cross-validated error of predicted R^2 (days and divergence uses both divergences). **b**, Closed loop, public MINDFUL data: Pearson correlation of angular error with elapsed days and each divergence (raw and partial given days), and leave-one-day-out mean absolute error (degrees).

Predictor	Raw ρ	Partial ρ
Elapsed days (log)	−0.67	–
Neural divergence	−0.34	−0.16
Output divergence	−0.22	−0.16
Sum	−0.36	−0.17
Cross-validated error of predicted R^2 (mean absolute error)		
Days only	0.085	
Divergence only	0.114	
Days and divergence	0.090	

b

				Neural		Output		Error (°)			
	Days	Windows	r(days)	raw	partial	raw	partial	days	+neural	+output	output
T11	15	145	0.84	0.82	0.37	0.81	0.73	17.1	17.5	11.9	12.8
T5	6	73	0.17	0.50	0.49	0.79	0.78	32.9	38.3	20.8	15.0

Supplementary Table 5 Channel controls and electrode link. For 40 training sessions per individual (22 for monkey M and 33 for T7; tuned penalty), median retention of the renormalized fixed decoder, after re-weighting the correctly assigned decoder, after re-weighting a decoder with scrambled channels, and the same two with non-negative weights. Electrode link (humans): median over training sessions of the Spearman correlation between fitted channel weights and the change in each electrode's log threshold-crossing rate or log impedance between the two days, with the share of sessions with a positive correlation.

			Signed weights		Non-negative weights		Electrode link	
Individual	Pairs	No labels	Correct	Scrambled	Correct	Scrambled	Firing rate	Impedance
Monkey N	134	0.39	0.46	−0.19	0.49	−0.17	–	–
Monkey C	84	0.42	0.62	0.02	0.59	0.01	–	–
Monkey M	43	0.49	0.68	0.01	0.67	−0.00	–	–
Human T6	113	0.61	0.89	0.65	0.89	0.36	0.05 (67%)	−0.02 (46%)
Human T5	112	0.13	0.83	0.77	0.72	0.57	0.03 (77%)	−0.01 (33%)
Human T9	108	0.15	0.78	0.56	0.72	0.29	0.11 (92%)	−0.07 (14%)
Human T7	73	0.09	0.80	0.63	0.67	0.38	0.04 (68%)	−0.06 (21%)
Human T10	92	0.23	0.72	0.57	0.65	0.35	0.06 (87%)	−0.03 (36%)
Human T8	117	−0.11	0.52	0.33	0.38	0.10	0.07 (90%)	−0.06 (18%)

Supplementary Table 6 Recurrent network versus linear decoder. For 25 training sessions per individual: median reference R^2, retention of the renormalized fixed decoder, and share of the loss recovered by re-weighting each channel.

		Reference R^2		Retention		Loss recovered	
Gap (days)	Pairs	Linear	Network	Linear	Network	Linear	Network
Monkey N							
1	8	0.28	0.40	0.77	0.77	−0.07	0.63
7	14	0.30	0.40	0.67	0.73	−0.24	0.52
30	19	0.28	0.36	0.41	0.48	−0.04	0.40
120	23	0.29	0.40	0.19	0.19	0.13	0.48
480	17	0.28	0.38	−0.00	−0.07	0.12	0.36
Monkey C							
1	17	0.31	0.66	0.71	0.81	0.41	0.68
7	10	0.35	0.74	0.59	0.73	0.38	0.52
30	4	0.50	0.87	0.23	0.15	0.35	0.45
120	6	0.25	0.71	0.00	0.04	0.41	0.57
480	12	0.14	0.32	−0.13	−0.39	0.22	0.45
Monkey M							
1	13	0.36	0.76	0.80	0.70	0.28	0.69
7	6	0.27	0.73	0.62	0.53	0.22	0.62
30	6	0.43	0.68	0.57	0.48	0.34	0.69
120	9	0.25	0.58	0.09	−0.50	0.35	0.64
480	9	0.35	0.83	0.04	−0.33	0.19	0.48
Human T6							
1	4	0.27	0.19	0.68	0.08	0.33	0.85
7	12	0.24	0.22	0.65	0.64	0.72	1.01
30	15	0.25	0.25	0.74	0.55	0.66	0.42
120	18	0.26	0.27	0.58	0.51	0.65	0.92
480	16	0.32	0.29	0.40	0.49	0.75	0.92
Human T5							
1	3	0.55	0.52	0.83	0.90	0.77	0.39
7	8	0.55	0.52	0.49	0.22	0.83	0.69
30	12	0.50	0.45	0.25	0.06	0.68	0.76
120	20	0.53	0.49	−0.17	−0.62	0.72	0.71
480	23	0.53	0.49	−0.27	−0.53	0.82	0.85
Human T9							
1	3	0.35	0.26	−0.43	−0.46	0.68	0.83
7	10	0.27	0.24	0.36	0.07	0.84	0.83
30	15	0.38	0.33	0.46	0.24	0.73	0.95
120	15	0.39	0.25	−0.02	−0.28	0.73	0.99
480	10	0.36	0.35	−0.33	−1.08	0.74	0.82
Human T7							
1	3	0.36	0.45	0.66	0.35	0.65	0.46
7	5	0.34	0.31	0.56	0.56	0.83	0.90
30	10	0.35	0.29	0.41	0.44	0.81	0.96
120	17	0.40	0.39	0.05	−0.04	0.80	0.82
480	10	0.57	0.57	−0.00	−0.34	0.80	0.84
Human T10							
1	7	0.35	0.28	0.83	0.48	0.72	0.83
7	12	0.31	0.26	0.37	−0.09	0.64	0.87
30	13	0.31	0.28	0.33	−0.34	0.71	0.81
120	10	0.27	0.26	−0.49	−1.39	0.58	0.40
Human T8							
1	2	0.36	0.33	0.28	−0.12	0.15	0.62
7	10	0.32	0.26	0.07	−1.07	0.08	0.27
30	12	0.40	0.37	0.22	−0.31	0.45	0.63
120	19	0.27	0.28	−0.21	−1.01	0.44	0.38
480	16	0.50	0.41	−0.23	−0.73	0.75	0.81

Supplementary Table 7 Matched-output control in the reaching monkeys (tuned penalty). Median retention of the renormalized fixed decoder and after re-weighting, and the share of the loss recovered by re-weighting (%), when decoding hand kinematics (four outputs) or hand-movement direction only (a two-dimensional unit vector, as for the human participants).

Individual	Target	Gap (days)	Renormalized	Re-weighted	Recovered (%)
Monkey C	kinematics	1	0.73	0.81	37
		7	0.62	0.73	38
		30	0.23	0.48	30
		120	0.04	0.35	36
		480	−0.12	0.10	22
	direction	1	0.83	0.90	48
		7	0.74	0.85	46
		30	0.37	0.71	60
		120	0.07	0.46	45
		480	0.03	0.38	36
Monkey M	kinematics	1	0.80	0.83	28
		7	0.62	0.70	22
		30	0.57	0.71	34
		120	0.09	0.34	35
		480	0.04	0.21	19
	direction	1	0.81	0.89	48
		7	0.73	0.87	36
		30	0.66	0.84	47
		120	0.34	0.52	30
		480	0.19	0.56	35

Supplementary Table 8 Angles between decoders fitted on different days (degrees, medians over pairs). Global, angle that allows one global scale; per channel, angle after fitting one scale per channel; permuted, the per-channel angle after randomly permuting the channels of the earlier decoder. Allowing one scale per channel reduced the angle by 8–37° in every individual, but permuted decoders given the same freedom reached 47–54° in the participants, and at the longest gap the per-channel angle was only 1–10° below that baseline. Angles that allow only a global scale can overstate rotation over days to months, but they say little about preserved tuning over years.

Individual	Gap (days)	Global	Per channel	Permuted
Monkey N	1	44	29	47
	7	50	33	49
	30	60	37	49
	120	65	40	48
	480	74	42	49
Monkey C	1	43	35	80
	7	52	40	75
	30	72	58	70
	120	75	56	78
	480	85	74	79
Monkey M	1	41	31	73
	7	49	39	79
	30	61	47	72
	120	72	53	70
	480	81	64	73
Human T6	1	47	31	53
	7	51	34	52
	30	53	35	53
	120	59	39	52
	480	64	42	52
Human T5	1	57	38	54
	7	65	43	53
	30	71	45	53
	120	79	48	54
	480	84	50	53
Human T9	1	52	33	49
	7	55	35	49
	30	59	37	49
	120	68	41	49
	480	77	44	49
Human T7	1	59	35	49
	7	59	36	52
	30	66	40	49
	120	72	43	48
	480	82	47	48
Human T10	1	49	34	50
	7	57	36	49
	30	61	39	49
	120	69	43	48
Human T8	1	47	28	47
	7	56	32	48
	30	62	36	49
	120	69	39	48
	480	82	45	48

Supplementary Table 9 Regularization by the 1% rule. The rule chose the largest ridge penalty whose median same-day R^2 was within 1% of the best. Columns give median R^2 with $\alpha = 0.1$ (the smallest penalty tested, for every individual) and with the chosen α, the median paired gain with 95% cluster-bootstrap confidence interval, and the share of training sessions whose mean gain was positive. The pair rule is not applied here, because penalties are compared on absolute R^2. †Interval not shown: a pair whose base decoder failed (for example R^2 near -300) makes the bootstrap interval meaningless. Monkeys here; the human participants follow on the next page.

Individual	Chosen α	Gap (days)	Pairs	R^2, $\alpha = 0.1$	R^2, rule	Gain [95% CI]	Share positive
Monkey N	10^4	same day	40	0.30	0.30	0.003 [0.001, 0.006]	0.72
Monkey N	10^4	1	11	0.27	0.28	0.021 [0.008, 0.035]	1.00
Monkey N	10^4	7	24	0.14	0.18	0.035 [0.026, 0.041]	0.92
Monkey N	10^4	30	28	0.10	0.14	0.042 [0.038, 0.052]	1.00
Monkey N	10^4	120	35	0.04	0.10	0.056 [0.050, 0.083]	0.97
Monkey N	10^4	480	28	−0.10	−0.00	0.090 [0.082, 0.112]	1.00
Monkey C	10^4	same day	40	0.21	0.22	0.002 [0.001, 0.003]	0.72
Monkey C	10^4	1	26	0.16	0.18	0.008 [0.005, 0.017]	0.88
Monkey C	10^4	7	22	0.20	0.21	0.012 [0.006, 0.019]	0.77
Monkey C	10^4	30	9	0.10	0.10	0.001 [−0.002, 0.011]	0.67
Monkey C	10^4	120	12	−0.01	0.01	0.031 [0.016, 0.045]	1.00
Monkey C	10^4	480	25	−0.10	−0.07	0.032 [0.014, 0.036]	1.00
Monkey M	10^3	same day	22	0.35	0.35	0.000 [0.000, 0.001]	0.73
Monkey M	10^3	1	14	0.26	0.26	0.001 [0.000, 0.001]	0.79
Monkey M	10^3	7	6	0.16	0.17	0.000 [−0.000, 0.002]	0.67
Monkey M	10^3	30	6	0.25	0.25	0.001 [0.000, 0.012]	1.00
Monkey M	10^3	120	9	0.01	0.01	0.001 [0.000, 0.002]	1.00
Monkey M	10^3	480	9	0.01	0.01	0.001 [0.000, 0.001]	1.00

Supplementary Table 9 (continued) Regularization by the 1% rule, human participants. Columns as on the previous page.

Individual	Chosen α	Gap (days)	Pairs	R^2, $\alpha = 0.1$	R^2, rule	Gain [95% CI]	Share positive
Human T6	10^4	same day	40	0.22	0.24	0.008 [0.003, 0.013]	0.95
Human T6	10^4	1	6	0.13	0.17	0.011 [0.002, 0.075]	1.00
Human T6	10^4	7	23	0.10	0.14	0.025 [0.016, 0.032]	1.00
Human T6	10^4	30	32	0.16	0.17	0.016 [0.013, 0.030]	0.97
Human T6	10^4	120	38	0.13	0.17	0.018 [0.009, 0.030]	0.95
Human T6	10^4	480	27	0.11	0.12	0.017 [0.005, 0.027]	0.89
Human T5	10^4	same day	40	0.51	0.52	0.008 [0.004, 0.014]	0.97
Human T5	10^4	1	13	0.37	0.45	0.038 [0.013, 0.073]	1.00
Human T5	10^4	7	13	0.26	0.26	0.018 [0.005, 0.053]	1.00
Human T5	10^4	30	25	−0.08	0.10	0.040 [0.012, 0.069]	1.00
Human T5	10^4	120	33	−0.02	0.01	0.040 [0.018, 0.053]	0.97
Human T5	10^4	480	40	−0.10	−0.03	0.029 [0.015, 0.060]	0.97
Human T9	3×10^4	same day	40	0.30	0.31	0.007 [0.000, 0.011]	0.65
Human T9	3×10^4	1	12	0.16	0.18	0.029 [0.013, 0.058]	0.92
Human T9	3×10^4	7	22	0.06	0.12	0.031 [0.012, 0.077]	0.82
Human T9	3×10^4	30	39	0.08	0.11	0.037 [0.027, 0.051]	0.87
Human T9	3×10^4	120	40	0.02	0.07	0.041 [0.029, 0.064]	0.90
Human T9	3×10^4	480	20	−0.17	−0.11	0.043 [0.035, 0.077]	1.00

Supplementary Table 9 (continued) Regularization by the 1% rule, human participants. Columns as on the first page.

Individual	Chosen α	Gap (days)	Pairs	R^2, $\alpha = 0.1$	R^2, rule	Gain [95% CI]	Share positive
Human T7	3×10^4	same day	33	0.24	0.26	0.007 [0.004, 0.024]	0.79
Human T7	3×10^4	1	5	0.13	0.14	0.016 (interval unstable†)	0.80
Human T7	3×10^4	7	10	0.08	0.10	0.039 [0.011, 0.170]	0.90
Human T7	3×10^4	30	24	−0.05	0.05	0.051 [0.023, 0.153]	1.00
Human T7	3×10^4	120	33	−0.14	−0.02	0.073 [0.044, 0.117]	0.91
Human T7	3×10^4	480	12	−0.00	0.04	0.028 [0.017, 0.094]	1.00
Human T10	3×10^4	same day	40	0.28	0.28	0.001 [−0.002, 0.003]	0.53
Human T10	3×10^4	1	15	0.04	0.08	0.021 [0.016, 0.041]	0.93
Human T10	3×10^4	7	20	0.09	0.10	0.016 [0.008, 0.026]	0.85
Human T10	3×10^4	30	34	0.00	0.05	0.042 [0.023, 0.052]	0.94
Human T10	3×10^4	120	32	−0.09	−0.05	0.031 [0.026, 0.045]	1.00
Human T8	3×10^4	same day	40	0.30	0.30	0.005 [0.001, 0.012]	0.72
Human T8	3×10^4	1	6	0.05	0.14	0.090 [0.012, 0.536]	1.00
Human T8	3×10^4	7	18	−0.03	0.04	0.060 [0.028, 0.154]	1.00
Human T8	3×10^4	30	25	−0.07	−0.00	0.037 [0.018, 0.052]	0.96
Human T8	3×10^4	120	36	−0.03	0.02	0.041 [0.027, 0.067]	0.97
Human T8	3×10^4	480	35	−0.17	−0.09	0.050 [0.038, 0.076]	1.00

a

Supplementary Table 10 Recalibration with few labeled trials. a, Monkey N (86 pairs, default penalty $\alpha = 0.1$): median retention without labels and, for each number of labeled trials, ridge to prior / fresh decoder / re-weighting (regularization by cross-validation). **b**, Human participants (tuned penalty; 30 training sessions each, all gaps; only pairs whose later day had at least 300 labeled training trials, so that every budget uses the same pairs): ridge to prior / fresh decoder / re-weighting, as in a. **c**, Harmful recalibrations in monkey N (default penalty; more than 0.01 R^2 below the renormalized fixed decoder) with the shrinkage chosen by cross-validation or fixed at the value best on other training sessions, with P from a two-sided Wilcoxon signed-rank test over training sessions (share of each session's pairs made worse) and Holm-adjusted over the five budgets, and median retention for both and for the best value in hindsight.

Gap (days)	Pairs	No labels	$n = 10$	$n = 20$	$n = 50$	$n = 100$	$n = 300$
1	7	0.64	0.47 / −0.18 / 0.02	0.63 / 0.22 / 0.50	0.75 / 0.55 / 0.64	0.83 / 0.70 / 0.69	1.03 / 1.02 / 0.73
7	18	0.48	0.51 / 0.17 / 0.38	0.58 / 0.39 / 0.43	0.77 / 0.68 / 0.54	0.89 / 0.81 / 0.61	1.02 / 1.01 / 0.54
30	20	0.15	0.21 / −0.07 / 0.05	0.27 / 0.25 / 0.09	0.60 / 0.61 / 0.25	0.79 / 0.80 / 0.29	1.02 / 1.02 / 0.27
120	26	0.13	0.31 / 0.24 / −0.02	0.44 / 0.43 / 0.06	0.67 / 0.66 / 0.24	0.82 / 0.83 / 0.19	1.01 / 1.01 / 0.30
480	15	−0.21	0.19 / 0.19 / −0.27	0.35 / 0.42 / −0.07	0.60 / 0.64 / 0.16	0.82 / 0.83 / 0.09	1.01 / 1.01 / 0.27

b

Participant	Pairs	No labels	$n = 10$	$n = 20$	$n = 50$	$n = 100$	$n = 300$
Human T6	41	0.56	0.63 / 0.36 / 0.54	0.68 / 0.47 / 0.53	0.72 / 0.68 / 0.64	0.75 / 0.76 / 0.61	0.90 / 0.92 / 0.84
Human T5	57	0.21	0.53 / 0.35 / 0.22	0.60 / 0.57 / 0.43	0.70 / 0.72 / 0.58	0.76 / 0.80 / 0.67	0.89 / 0.91 / 0.81
Human T9	24	0.14	0.14 / 0.12 / 0.26	0.27 / 0.55 / 0.27	0.36 / 0.79 / 0.46	0.50 / 0.81 / 0.65	0.93 / 0.98 / 0.75
Human T7	24	0.09	0.54 / 0.52 / 0.48	0.65 / 0.83 / 0.63	0.79 / 0.93 / 0.70	0.82 / 0.92 / 0.69	0.99 / 0.99 / 0.78
Human T10	40	0.29	0.25 / 0.11 / 0.22	0.36 / 0.34 / 0.23	0.49 / 0.60 / 0.44	0.60 / 0.74 / 0.59	0.89 / 0.94 / 0.72
Human T8	59	0.04	0.21 / 0.10 / −0.07	0.32 / 0.16 / 0.06	0.36 / 0.41 / 0.18	0.44 / 0.46 / 0.29	0.74 / 0.71 / 0.54

c

		Harmful (%)				Median retention		
Trials	Pairs	Cross-validation	Fixed	P	P (Holm)	Cross-validation	Fixed	Hindsight
10	86	10.5	1.2	0.027	0.136	0.31	0.37	0.40
20	86	8.1	0.0	0.042	0.169	0.43	0.44	0.48
50	86	1.2	0.0	0.317	0.952	0.64	0.63	0.68
100	86	0.0	1.2	0.317	0.952	0.82	0.81	0.83
300	86	0.0	0.0	1.0	1.0	1.02	1.01	1.02

a

Supplementary Table 11 Electrode statistics. a, Silence and recovery in initially active human electrodes under different definitions, observed and after shuffling each electrode's session order (mean of 200 shuffles; mean number of shuffled silent episodes in parentheses). **b**, Per array: yield is the percentage of electrodes with a threshold-crossing rate of at least 2 Hz (median of the first / last three sessions); h_{off} and h_{on} are switching rates per 1,000 days; Rev./sil., recovered and silenced initially active electrodes; Kurt., excess kurtosis of session-to-session changes in log rate after removing session-wide changes; Imp. ρ, Spearman correlation of the array median impedance with day (monkey N: median over channels of log impedance with day); Edge − int., median edge minus median interior slope of log rate per year among initially active electrodes, and Adjusted, the same after adjusting for each electrode's initial rate (the values tested in the main text with the array as the unit, and shown in Fig. 7d; arrays with fewer than 10 initially active electrodes, –). Monkey N's edge value pools both arrays and all electrodes.

Definition	Silent episodes	Recovered	Recovered (%)	Shuffled (%)
At least 3 sessions below 2 Hz (main definition)	730	584	80	78 (601)
Threshold −3.5 times noise	8	7	88	–
Threshold −5.5 times noise	646	475	74	–
Fixed threshold in microvolts per electrode	840	601	72	74 (652)
Array-wide shifts removed	747	575	77	–
Mean waveform trough at least 30 μV	736	523	71	–
Silence below 0.5 Hz	607	430	71	70 (434)
Silence lasting at least 30 days	168	96	57	44 (79)

b

Supplementary Table 11 (continued) Electrode statistics. Panel b, as described on the previous page.

Part.	Array	Sess.	Span (d)	Yield (%)	h_{off}	h_{on}	Rev./sil.	Kurt.	Imp. ρ	Edge − int.	Adjusted
A1	single	48	252	81 / 55	16.3	27.1	44/68	2.9	−0.27	0.78	−0.35
S1	single	40	337	3 / 1	46.6	10.7	1/4	0.7	–	–	–
S2	single	32	458	1 / 6	24.2	10.9	0/1	–	–	–	–
S3	single	207	1,954	84 / 10	14.9	3.3	68/88	3.9	−0.91	−0.11	−0.11
T1	single	41	277	65 / 86	2.6	27.1	11/17	7.8	−0.93	−0.42	−0.24
T10	MFG	104	329	96 / 86	8.6	54.1	56/59	9.8	−0.70	0.33	0.60
T10	dPCG	104	329	21 / 44	7.7	9.1	26/31	3.2	−0.21	−0.19	−0.02
T11	lateral	292	1,677	20 / 10	7.0	1.5	16/24	5.5	−0.95	−0.09	−0.11
T11	medial	292	1,677	38 / 64	3.7	6.3	35/35	4.1	−0.99	−0.12	−0.08
T2	single	176	914	57 / 14	22.4	3.4	50/67	4.1	−0.98	0.04	0.40
T3	single	87	399	0 / 2	34.3	4.0	0/0	–	−0.98	–	–
T5	lateral	438	2,665	60 / 30	2.4	1.2	41/59	4.4	−0.94	−0.05	−0.05
T5	medial	438	2,665	51 / 40	2.9	2.7	54/59	2.4	−0.90	−0.08	−0.02
T6	single	190	1,184	16 / 2	11.0	3.5	12/22	2.7	−0.97	−0.20	0.01
T7	lateral	84	521	24 / 31	20.1	21.7	23/28	1.9	−0.83	−0.62	−0.24
T7	medial	84	521	36 / 49	8.7	11.8	34/42	1.5	−0.91	−0.51	−0.28
T8	lateral	329	1,061	11 / 38	1.1	1.2	12/12	6.7	−0.92	−0.80	−0.64
T8	medial	329	1,061	55 / 41	5.2	2.9	64/70	3.0	−0.92	0.33	0.12
T9	lateral	221	731	14 / 16	18.7	7.5	13/17	4.3	0.26	−0.96	−0.65
T9	medial	221	731	21 / 44	13.1	6.6	24/27	2.1	−0.96	−0.01	0.11
Monkey N	both	303	1,242	38 / 19	3.1	0.8	23/31	10.2	−0.88	−0.16	–

Supplementary Table 12 Simulator parameters. Values calibrated on monkey N sessions before day 700 (early) and from day 700 onwards (late). Switching rates were measured from electrode activity in each period. Fixed parameters: local mixing share $\beta = 0.5$, turnover concentration 5, 20 latent dimensions and silent-signal fraction $\kappa = 0.37$.

Parameter	Unit	Early	Late
Session-to-session mixing s.d.	1	0.959	0.768
Slow mixing s.d. (asymptote)	1	1.478	1.317
Slow mixing time constant	days	15	12
Turnover fraction, next session	1	0.042	0.032
Turnover fraction, asymptote	1	0.120	0.127
Turnover time constant	days	360	527
Active to silent rate	per day	0.0038	0.0034
Silent to active rate	per day	0.0018	0.0017
Calibration loss		0.074	0.063

Supplementary Table 13 Training on simulated futures. Median R^2 on real later sessions of monkey N for ridge decoders trained on sessions after day 700 (simulator calibrated before day 700). Base, $\alpha = 0.1$; stronger ridge, $\alpha = 10^4$; perturbation, random channel dropout and gain noise; simulator, eight simulated futures; simulator and ridge, both; several days, data from the three previous sessions; reference, decoder trained on the test day.

Gap (d)	Pairs	Base	Stronger ridge	Perturb.	Simulator	Sim. + ridge	Several days	Reference
1	9	0.21	0.22	0.20	0.20	0.20	0.23	0.29
7	20	0.18	0.21	0.18	0.19	0.20	0.23	0.34
30	18	0.08	0.13	0.08	0.09	0.10	0.15	0.28
120	18	0.07	0.12	0.07	0.07	0.08	0.14	0.30
240	12	0.05	0.10	0.06	0.08	0.08	0.11	0.29

Supplementary Table 14 BrainGate decoding participants and the inclusion rule. Arrays (channels), sessions and span of the closed-loop cursor sessions; median trials (movement periods) per session and median movement duration; median R^2 of the reference decoder across session pairs (linear decoder, default penalty). Participants entered the decay and correction analyses when this median exceeded 0.1 (Methods). For T2, neither block-wise normalization nor threshold crossings raised same-day R^2 above about 0.1 except on its earliest session.

Participant	Arrays (ch.)	Sessions	Span (yr)	Trials/session	Movement (s)	Ref. R^2	Included
T6	1 (96)	124	3.1	368	1.6	0.25	yes
T5	2 (192)	92	7.2	552	1.7	0.50	yes
T9	2 (192)	117	2.0	275	3.9	0.33	yes
T7	2 (192)	35	1.4	290	3.8	0.32	yes
T10	1 (96)	57	0.9	403	3.2	0.26	yes
T8	2 (192)	58	2.7	722	2.8	0.34	yes
T2	1 (96)	42	2.0	184	15.6	−0.02	no
T3	1 (96)	8	1.0	158	15.6	−0.05	no
T11	2 (192)	196	4.6	215	3.3	−16.19	no

Supplementary Table 15 Benchmark baselines. Valid pairs per individual and median retention of three reference methods after 1, 30 and 480 days: the renormalized fixed decoder (no labels), re-weighting each channel and a refit shrunk toward the old decoder (both with 300 labeled trials). Linear decoder, tuned penalty. T10 had no 480-day pairs. Per-pair scores and bootstrap intervals are released with the benchmark.

		Renormalized			Re-weighted (300)			Refit to prior (300)		
Individual	Pairs	1 d	30 d	480 d	1 d	30 d	480 d	1 d	30 d	480 d
Monkey N	259	0.79	0.41	0.01	0.80	0.45	0.28	1.02	1.00	1.01
Monkey C	130	0.73	0.23	−0.12	0.81	0.48	0.10	1.00	1.00	1.00
Monkey M	43	0.80	0.57	0.04	0.83	0.71	0.21	1.01	1.01	1.02
Human T6	333	0.79	0.76	0.35	0.88	0.92	0.82	0.92	0.98	0.96
Human T5	247	0.73	0.28	−0.17	0.92	0.82	0.82	0.94	0.94	0.95
Human T9	320	0.60	0.31	−0.38	0.84	0.82	0.70	0.97	0.98	0.97
Human T7	73	0.53	0.29	0.06	0.72	0.84	0.78	0.99	0.99	0.99
Human T10	124	0.58	0.18	–	0.85	0.73	–	0.96	0.95	–
Human T8	155	0.49	−0.03	−0.24	0.66	0.55	0.58	0.60	0.65	0.86

Supplementary Table 16 Pretrained foundation decoder (NDT3) against the ridge decoder. Both decoders received threshold crossings on the same 12 session pairs per individual (gaps of 1 to 480 days). Each cell gives the median R^2 of NDT3 / the ridge decoder, and in parentheses the number of sessions (same day) or pairs where NDT3 was higher. No update: NDT3 applied unchanged versus the renormalized ridge decoder. 50 and 300 trials: NDT3 fine-tuned with the first labeled trials of the new day versus the ridge decoder refit toward its previous weights. Self-supervised: median retention of NDT3 after continued training with its neural prediction loss only. T8's nearly silent first array was left out of NDT3's input by a rule written after a diagnostic on T8 and before its pairs were run (Methods); the ridge decoder used all channels.

Individual	Pairs	Same day	No update	50 trials	300 trials	Self-supervised
Monkey N	12	0.42 / 0.20 (12/12)	−0.06 / −0.04 (2/12)	0.16 / 0.08 (10/12)	0.41 / 0.18 (12/12)	−0.43
Human T5	12	0.34 / 0.45 (0/11)	−0.24 / −0.19 (3/12)	0.06 / 0.30 (2/12)	0.26 / 0.41 (1/12)	−0.16
Human T8	12	0.22 / 0.28 (4/12)	−0.09 / −0.10 (4/12)	−0.12 / 0.07 (0/12)	−0.04 / 0.18 (1/12)	−0.08

Supplementary Table 17 Decay summary. Retention of the renormalized fixed decoder (default penalty) after exactly one day (median, 95% confidence interval and number of one-day pairs; shown when there were at least three) and days until median retention fell to one half, with the default and the tuned ridge penalty. Intervals are from a cluster bootstrap over training sessions (2,000 resamples). <1, below one half at the first tested gap. Pairs, valid pairs with the default penalty.

Individual	Pairs	One-day retention	Half-life, default (days)	Half-life, tuned (days)
Monkey N	259	0.74 (0.62 to 0.83; 21)	8.1 (3.7–10)	16 (13–42)
Monkey C	130	0.70 (0.58 to 0.78; 27)	9.4 (8.0–24)	11 (8.8–34)
Monkey M	43	0.78 (0.69 to 0.85; 11)	34 (17–45)	37 (21–48)
Human T6	311	0.56 (−1.09 to 0.99; 4)	138 (<1–251)	206 (108–311)
Human T5	247	– (2)	3.2 (<1–8.8)	8.1 (<1–27)
Human T9	318	0.58 (0.33 to 0.76; 17)	1.8 (<1–4.0)	2.6 (<1–7.5)
Human T7	69	– (2)	<1 (<1–10)	7.5 (<1–18)
Human T10	117	0.56 (−0.01 to 0.84; 10)	1.8 (<1–10)	2.7 (<1–12)
Human T8	155	0.57 (0.00 to 0.89; 4)	<1 (<1–3.2)	<1 (<1–4.2)

Supplementary Table 18 Alignment ceiling. Share of the loss recovered at gaps of 120 and 480 days (percent, with 95% cluster-bootstrap intervals), relative to the renormalized fixed decoder (0%) and the reference decoder (100%), for the label-free Procrustes rotation of the dominant 16-dimensional subspace, the best alignment of that subspace fitted with all labeled training trials of the later day (ceiling), and re-weighting of each channel fitted on the same trials (tuned penalty). Last column: median over training sessions of the paired difference in R^2, re-weighting minus ceiling, over all gaps, with two-sided Wilcoxon signed-rank P; positive values mean per-channel re-weighting recovered more than the best linear alignment of the subspace. Pairs: all gaps; the shares use gaps of 120 and 480 days.

Individual	Pairs	Label-free rotation	Ceiling	Re-weighting	Re-weighting − ceiling
Monkey N	259	−10 (−17 to −4)	48 (40 to 53)	15 (7 to 21)	−0.06 (2e-14)
Monkey C	130	−27 (−93 to −14)	77 (62 to 82)	46 (32 to 53)	−0.04 (8e-7)
Monkey M	43	−21 (−41 to −3)	85 (72 to 87)	22 (16 to 29)	−0.04 (5e-7)
Human T6	333	0 (−15 to 8)	70 (61 to 76)	75 (67 to 79)	0.01 (5e-4)
Human T5	247	−8 (−19 to −2)	77 (74 to 80)	86 (84 to 88)	0.04 (4e-15)
Human T9	320	−7 (−20 to 1)	61 (54 to 64)	78 (72 to 80)	0.05 (2e-17)
Human T7	73	−18 (−32 to 0)	63 (44 to 72)	84 (70 to 88)	0.06 (2e-8)
Human T10	124	0 (−11 to 19)	49 (42 to 73)	75 (64 to 92)	0.04 (2e-7)
Human T8	155	−13 (−25 to −5)	42 (34 to 47)	76 (70 to 82)	0.14 (1e-10)